

ELI — Adaptive Inference Governor

ELI v2.4.76

September 2026

Contents

Adaptive Inference Governor — Status and Remaining Plan	1
0. The one-sentence problem	1
1. What ELI already adapts to (do not rebuild)	1
2. The gaps and where each stands	2
3. Design for what remains: the Adaptive Inference Governor	3
4. Cross-machine behaviour to aim for	4
5. Phases	4
6. Verification	5
7. Risks and non-goals	5
8. Mixture-of-experts models on a small GPU	5
9. The startup dialog's own model loader	5

Adaptive Inference Governor — Status and Remaining Plan

Updated for v2.4.72. Most of the gaps this plan identified in June are closed; the two that need a governor (throughput-aware think decision and per-call wall-time projection) are still open. Status per gap below; no governor code exists yet.

Scope: make ELI's reasoning and inference *policy* (think or no-think, token budget, pass count, recovery) a pure function of the runtime environment and the task, so one redistributed binary behaves correctly on any machine with any model and no per-machine tuning.

Origin: a 2026-06-17 session in which EXPLAIN_COGNITION_RUNTIME burned a 273 s empty compact-synthesis pass, then two 539 s empty standard-synthesis passes, before a two-stage chain-of-thought fallback (470 s and 502 s): about 38 minutes for one question. The working final answer was 4,071 characters (about 1,018 tokens) in 502 s, roughly 2 tokens per second.

0. The one-sentence problem

ELI adapts its inference to **memory** (VRAM, context length) and **model size**, but not to **measured compute speed** (tokens per second). On a slow machine the same prompt that works on the development box can produce multi-minute turns.

1. What ELI already adapts to (do not rebuild)

Capability	Where	What it adapts to
Model-agnostic resolution (no baked model)	gguf_inference.get_model_path	settings, environment
Chat family and template	cognition/model_identity.py, model_output_tokens.py	GGUF metadata (template, architecture), not the file name
Thinking-channel detection	gguf_inference._is_thinking_model → model_identity.is_thinking_model	the embedded chat template, then architecture; no name list
Trained context length	core/startup_hardware_optimizer, ELI_MODEL_TRAINING	ELI_MODEL_TRAINING, then real GGUF context_length, else 8192; no filename table
Free-VRAM-aware sizing, one fit calculation	core/hardware_profile.py (smart_fit_config, unified_fit_config)	free VRAM, RAM budget, fit priority
KV-cache and compute-buffer reserve	hardware_profile.py (_kv_cache_mb, _compute_graph_reserve_mb)	context, layers, batch
Graceful GPU-layer fallback	boot optimiser; runtime_snapshot.json	load success
Per-mode token budget	adaptive_load_report cognition/reasoning_modes.py	n_ctx (20 % of the window for quick, 30 % otherwise), prompt pressure, complexity, mode
Size tier → budget scale	detect_tier, tier_scale	model size
Live decode-speed EMA	model_tier.record_speed, measured_tok_s (fed from the non-stream generation path in gguf_inference)	measured tokens per second
Speed-capped pass count	model_tier.speed_passes	measured tokens per second (thresholds ELI_SLOW_TPS=5, ELI_FAST_TPS=15)
No-think on utility calls	gguf_inference._no_think_prefill, force_no_think()	structured output, small budgets (< 1024), quick mode
Redistribution-aware settings	core/runtime_settings.py	machine moves

The speed EMA is consumed in exactly one place, speed_passes(). Nothing in the think decision, the token budget or mode selection reads it. speed_passes() reduces how many full generations a mode runs and never caps output length.

2. The gaps and where each stands

Gap	Was	Status
A thinking detection was a filename allowlist	_is_thinking_model matched qwen3, deepseek-r1, ...	Closed. Detection reads the chat template, then the architecture.
B the think decision and token budget ignored throughput	no tok/s term in _no_think_prefill or the budget	Open. The inputs are structured, max_tokens < 1024, the quick mode, the force_no_think() scope and ELI_MODEL_THINK.

Gap	Was	Status
C no wall-time projection	nothing computes “at the live EMA this call takes about N seconds”	Open.
D empty output was not self-correcting	an unterminated <think> returned empty, then a full-price retry	Partly closed. An empty non-stream response is retried once with <code>force_no_think</code> in the engine (<code>_get_chat_response</code>), and an empty stream is retried once with <code>no-think</code> in <code>gguf_inference</code> . There is still no reserved answer budget.
E context sizing undershot the trained context	a filename table returned 32768 for <code>qwen3</code>	Closed. GGUF metadata is authoritative.
F <code>force_no_think</code> was a point patch	one call site	Partly closed. It now wraps compact grounded synthesis, the chain-of-thought stage prose, the correction repair, and the empty-response and stream retries in <code>_get_chat_response</code> . The standard synthesis fallback (<code>_synthesize_answer</code>) and coding-agent fix generation for <code>SELF_IMPROVE</code> are not wrapped.
G redundant proactive churn	<code>persona_updater</code> rebuilt the same overlay every tick	Closed. <code>Persona</code> and <code>user-profile</code> overlays are debounced (120 s) with a content dirty-check.

Note on gap G: `persona_updater` is deterministic (database reads and regexes); the cost was redundant CPU and SQLite work, not inference-lock contention.

3. Design for what remains: the Adaptive Inference Governor

A single policy object that every model call routes its *decision* through. It consumes `reasoning_modes`, `model_tier` and `hardware_profile`; it does not replace them.

3.1 Inputs (all available at runtime)

- **Model capability:** `model_identity.is_thinking_model(template, architecture)`, cached per model.
- **Throughput:** `model_tier.measured_tok_s()`, with a cold-start default from `detect_tier()` until the first measurement.
- **Context:** `n_ctx` from the runtime snapshot and `n_ctx_train` from metadata.
- **Task class:** below.

3.2 Task classes

1. **STRUCTURED:** JSON, routing, plan graphs. Never thinks (already true via `structured=True`).
2. **GROUNDING_SYNTHESIS:** evidence already gathered; the call only phrases it (the `EXPLAIN_*` actions, RAG answer synthesis, fix-patch rendering). Thinking adds nothing; force no-think.
3. **OPEN_REASONING:** genuine multi-step reasoning (chain of thought, tree of thoughts, constitutional, self-consistency, open chat on a hard question). Thinking is allowed but bounded by throughput.

3.3 The decision (a pure function)

```
policy(task_class, model_caps, tok_s, n_ctx, prompt_tokens) -> {  
  think: bool,  
  think_budget: int,      # max tokens allowed inside <think>  
  answer_budget: int,     # reserved tokens after </think>  
  passes: int,            # via speed_passes()  
}
```

- A model that does not think: think=False always.
- STRUCTURED or GROUNDED_SYNTHESIS: think=False.
- OPEN_REASONING: at or above ELI_FAST_TPS think with the full budget; between slow and fast scale think_budget so the projected wall time is within the target; at or below ELI_SLOW_TPS turn thinking off, or hard-cap think_budget.
- Always reserve answer_budget, so a thinking call cannot exhaust its budget before it answers.

3.4 Integration points

- gguf_inference._no_think_prefill: consult the governor instead of only structured, < 1024 and the environment; keep force_no_think() as the primitive and ELI_MODEL_THINK as the hard user override.
- reasoning_modes: split the target tokens into think_budget and answer_budget.
- Engine synthesis sites: tag_compact_grounded_synthesis and _synthesize_answer as GROUNDED_SYNTHESIS and the reasoning-mode runners as OPEN_REASONING.
- Feed the speed EMA from the stream path as well, so quick-mode chat updates it.

3.5 Configuration

ELI_SLOW_TPS and ELI_FAST_TPS exist (5 and 15). A soft per-turn wall-time target (ELI_TURN_LATENCY_TARGET_S) is proposed and does not exist yet. ELI_MODEL_THINK is the hard override. Nothing machine-specific is persisted.

4. Cross-machine behaviour to aim for

Machine and model	tok/s	grounded actions (EXPLAIN_*)	hard chat
8 GB GPU with a 30B MoE	about 2	no-think, one pass	bounded think, within the target, never empty
24 GB GPU with a 7B	about 40	no-think	full think, full budget
CPU-only with a 7B	about 3	no-think	think off or hard-capped, with an answer reserve
any machine with a non-thinking model	n/a	normal answer	normal answer
any machine with an unfamiliar reasoning model	any	detected through the template	handled

5. Phases

1. **Governor core** for grounded synthesis: the policy function, no-think wired through _no_think_prefill, coding-agent fix rendering tagged.
2. **Throughput-bounded open reasoning** with an answer reserve: split budgets in reasoning_modes.

3. **EMA from streaming** so quick mode updates the measured speed.

Each phase is independently shippable.

6. Verification

- Unit: `policy(...)` truth table over task class $\times$ throughput tier $\times$ model-thinks.
- No-model engine tests in the style of `tests/test_*_no_gguf*.py`.
- Behavioural: `EXPLAIN_MEMORY_RUNTIME` and `EXPLAIN_COGNITION_RUNTIME` synthesis is non-empty on the first pass and runs no-think.
- Speed simulation: inject a fake `measured_tok_s()` and assert the think flag and budget flip.
- Existing guards that must stay green: `tests/test_gguf_think_stop_collision.py`, `tests/test_reasoning_mode_collision.py`, `tests/test_cot_think_budget.py`.

7. Risks and non-goals

- **Over-suppressing thinking** on hard turns at mid speed: only `GROUNDED_SYNTHESIS` is unconditionally no-think; `OPEN_REASONING` keeps thinking with a budget above slow.
- **Non-goals:** changing model selection, the agent set, persona content or the grounding contract; removing `ELI_MODEL_THINK`.

8. Mixture-of-experts models on a small GPU

`eli/core/moe_offload.py`. A mixture-of-experts model uses a few experts per token, and the experts are most of the file. Splitting such a model by layers puts mostly idle experts on the GPU. With the offload on, every layer’s attention and shared weights go to the GPU and the expert tensors stay in system memory (llama.cpp’s tensor buffer override, applied to the model parameters while the model is built, because llama-cpp-python does not expose it).

- Measured on an RTX 2060 SUPER (8 GB), Qwen3.6-35B-A3B Q4_K_M, 1,500-token prompt: 8 GPU layers gave 71 tok/s prompt and 8.3 tok/s generation; all layers with experts in RAM gave 125 tok/s and 21.8 tok/s. Load time fell from 63 s to 5 s. Through `load_model` the same model went from 10.7 to 20.8 tok/s.
- When: the model is MoE (`GGUFModelProfile.is_moe`), it would not fit in free VRAM, and available RAM covers about three quarters of the expert weights. Otherwise the ordinary layer split is used.
- Switch: the “Expert offload (MoE)” choice in the startup dialog, the setting `moe_expert_offload` or env `ELI_MOE_EXPERT_OFFLOAD = auto` (default), on, off. `ELI_MOE_EXPERT_FRACTION` adjusts the assumed expert share of the file (default 0.90).
- The live runtime snapshot carries `moe_expert_offload`.

9. The startup dialog’s own model loader

`eli/gui/eli_pro_audio_gui_v2_0.py` builds `Llama(...)` itself for the “Load selected model now” path — it does not go through `gguf_inference.load_model()`. The MoE expert-offload hook (section 8) is wired into that loader too now; before this it was saved by the toggle but never reached the actual load, so the setting did nothing when a model was loaded from the startup dialog.