ELI — Architecture

ELI v2.3.2

August 2026

Contents

Blueprint — ELI MKXI Architecture

The full structural/architectural map of ELI MKXI. Every claim here is grounded in the actual source tree (file paths are real and clickable). Where something is observed-at-runtime rather than read-from-code it is marked *(runtime)*.

> ELI is a **local-first, offline-by-default, model-agnostic** cognitive runtime + assistant GUI (and, as of 2026-06-28, a self-hosted web app — `api/server.py`, documented in `ELI_USER_MANUAL.md`). No cloud, no APIs on the inference path, no hardcoded model. ~160k LOC across 397 Python files (+ `api/server.py`); 216 capabilities (2026-08-06).

---

0. Design principles (the non-negotiables)

1. **Local-first / offline-by-default.** All outbound network is fail-closed at the socket boundary (`eli/core/netguard.py`); networked features opt in.

2. **Model-agnostic.** No model name/size is hardcoded on the inference path; ELI discovers and loads whatever GGUF exists. The model is swappable substrate, not a constant.
3. **Grounded / anti-confabulation.** A large dedicated subsystem exists purely to stop a local model stating things it can't back up (see §9).
4. **Emergent persona.** ELI has a distinctive voice; persona drift/banter is intended. Functional bugs are fixed; the personality is not scripted away.

---

1. Repository map (recursive LOC, role)

Per-package files / LOC measured 2026-07-28.

Package	LOC	Files	Role
`eli/runtime`	30.4k	94	Grounding spine, evidence, response/introspection surfaces, daemons, self-improvement
`eli/execution`	24.2k	15	Router + executor: intent → action → side effects (executor god-file)
`eli/gui`	22.7k	24	PySide6 desktop app, panels, startup/first-boot wizard, animated face
`eli/cognition`	15.4k	31	Agent bus, 12-stage orchestrator, inference, persona, reasoning, tone/emotion adaptor, emotional timeline
`eli/kernel`	14.7k	8	`CognitiveEngine` (the orchestrating core), scheduler, task bus
`eli/perception`	8.3k	23	Vision (VL), STT (whisper), TTS, **wake word**, **voice/tone**, OS control, gaze
`eli/core`	8.1k	28	netguard, paths, settings, hardware profile, model download, **full_control**
`eli/tools`	7.4k	29	Image engine, news, registry/capabilities, document tools

Package	LOC	Files	Role
`eli/memory`	7.4k	13	SQLite + FTS5 + FAISS + knowledge graph + working memory
`eli/planning`	4.0k	24	Proactive daemon, habit scheduler, job/proposal/attention queues
`eli/learning`	3.4k	12	LoRA self-training pipeline (Phi-3 base), dataset build/eval
`eli/coding`	2.1k	12	`CodeAgent` — plan→search→verify→repair coding pipeline
`eli/plugins`	2.0k	26	Plugin manager + bundled plugins
`eli/world`	1.6k	26	EliWorld event bus, local world bridge, avatar/ontology/face
`eli/integrations`	1.1k	9	Optional Ollama client + integration adapters (not on the default path)
`eli/utils`	1.0k	3	logging, shared helpers
`eli/contracts`	0.7k	3	typed pipeline contracts
`eli/system`	0.3k	2	system-level helpers
`eli/cli`	0.1k	2	headless REPL
total (listed)	**~155k**	**383**	full `eli/` tree: ~160k / 397 incl. small top-level pkgs

The four god-files (refactor targets — see §20)

File	LOC
`eli/execution/executor_enhanced.py`	14951
`eli/kernel/engine.py`	13841
`eli/gui/eli_pro_audio_gui_v2_0.py`	12100
`eli/execution/router_enhanced.py`	7621
`eli/runtime/deterministic_grounding_gate.py`	4301
`eli/memory/memory.py`	4734

---

2. Entry points & launch

Command	Target	Notes
`eli`, `eli-mkxi`, `eli-cli`	`eli.gui.app:main`	GUI (default)
`eli-gui`	`eli.gui.app:main`	GUI script
`python -m eli`	`eli/__main__.py`	dispatches to GUI or headless
`python -m eli --headless` / `-H`	`eli.cli.headless:run_headless`	terminal REPL, no Qt
`./eli.sh`	`python -m eli`	venv launcher
`bash install.sh`	—	venv + deps + clean config seed + verify

Boot path (GUI): `app.py` → (first-boot wizard if no model, `eli/gui/panels/startup.py`) → `eli_pro_audio_gui_v2_0.py:main` builds `EliMainWindow` → constructs the **`CognitiveEngine` singleton** → loads GGUF per `config/settings.json` → starts daemons.

---

3. The request lifecycle (the spine)

Everything funnels through **`CognitiveEngine.process()`** (`eli/kernel/engine.py`). Two paths diverge by reasoning mode:

```
user input
   |
   ▼
[Router]  eli/execution/router_enhanced.py :: route()
   |   priority pipeline → {action, args, confidence, meta.matched_by}
   ▼
CognitiveEngine.process(action, args, ...)
   |
   ├─ PHASE45 fast-path?  (deterministic OS/media/status/job actions)
   |      → execute_action() → return VERBATIM (no LLM)        [§10]
   |
   ├─ non-quick mode OR forced → Internal Orchestrator (full 12-stage)  [§7]
   |
   └─ quick mode (default) →
          AgentBus.dispatch()   [§6]  → bus_result {grounding, agents, ctx}
               |
               ├─ Grounding escalation hook  [§9]  (CHAT + low grounding + factual)
               |      → web tier / local tier / hedge → may RETURN here
               |
               ├─ self-contained action? → return executor result verbatim
               ├─ grounded control action? → grounded synthesis
               └─ CHAT → _build_enhanced_system() → broker.infer()/stream
                         → output_governor → response
```

Pipeline stages logged at runtime: `Stage 1 Intent → … → Stage 12 Confidence`. Quick mode short-circuits stages 2–4/HyDE; the orchestrator runs all 12.

---

4. Routing layer (`eli/execution`)

- **`router_enhanced.py`** — `route(text) -> {action, args, confidence, meta}`. Regex-first **priority pipeline** (“explicit priority pipeline installed”): ordered prepasses (web realtime lookup, media, file/PDF, memory, control…) then a core router, with an LLM-intent fallback (`cognition/llm_intent.py`). Holds module state like `_last_used_path` for referential follow-ups.

- **`execution_planner.py`** / `execution_intent_packets.py` — typed `ExecutionPlan` / `RouteDecision` artifacts the bus consumes.
- **`route_authority.py`, `route_contracts.py`, `tool_execution_authority.py`, `operator_policy.py`** — guardrails on what may be routed/executed.
- **`portable_intent_contract.py`, `router_plugin_intents.py`, `executor_plugin_handlers.py`, `media_runtime.py`, `operator_actions.py`** — plugin/media/OS intent wiring.

Key router behaviours (each is a fixed bug → guarded by `tools/eval/cases.yaml`): media controls require **whole-word + command-shape** (no substring “tab” in “metabolise”); bare “do a search” needs a subject; meta-questions stay CHAT; realtime facts (“release date”) → WEB_SEARCH.

5–6. Agent Bus & the agents (`eli/cognition/agent_bus.py`)

`AgentBus.dispatch(user_input, intent, …) -> DispatchResult`. Selects an agent set (`_select_agents_for_intent`, or broad fan-out), runs them over a dependency DAG (falls back to flat parallel), aggregates.

`DispatchResult` fields: `agent_results[]`, `action_result`, `memory_context`, `aggregated_confidence` (route+grounding blend), **`grounding_confidence`** (agent-evidence only — the honest “is this backed by anything” signal), `agents_used`, `confidence_label`, `orchestrator_plan`.

14 registered bus agents (`_ALL_AGENTS`):

Agent	Role
`BusMemoryAgent`	recall from SQLite/FTS/FAISS into context
`KnowledgeGraphAgent`	entities/relations from the KG
`SystemAgent`	runs executor actions (WEB_SEARCH, RUNTIME_AUDIT, …)
`OrchestratorAgent`	plan/coordinate grounded synthesis
`FileCodeAgent`	searches the codebase for code-grounded answers
`IntrospectionBusAgent`	live runtime/self introspection
`CapabilityAgent`	what ELI can do (capability manifest)
`ReflectionAgent`	session reflection/insights
`ProactiveAgent`	proactive patterns/signals
`HabitAgent`	learned habits
`SelfImprovementAgent`	self-improvement state
`FrontierAgent`	frontier reasoning hooks
`PluginAgent`	dispatches to registered plugins
`VoiceAgent`	STT/TTS engine state

+ the 15th: `CodeAgent` (`eli/coding/agent.py`) — a *separate* coding pipeline (plan → search → verify → repair, beam/DAG), invoked for CODE_SOLVE / GENERATE_SCRIPT, not a bus agent.

---

7. The Orchestrator — full 12-stage (`eli/cognition/orchestrator.py`)

Runs for non-quick modes (and forced cases). Stages (as logged):

1. **Intent Routing** → action
2. **Persona Lock** → persona/identity fixed for the turn
3. **HyDE** → hypothetical-doc query expansion (`cognition/hyde.py`)
4. **Planner** → which retrieval channels + limits for the mode 5/6/7. **Parallel Retrieval** → keyword + **FTS5** (conversation_turns) + **FAISS** vector + RAG + **KG**
5. **Hybrid Merge** → combine channel hits
6. **Rerank** → `cognition/reranker.py`

7. **Context Assembly** → assembled grounded context 10.5. **Persona Handoff** → persona brief built for generation
8. **LLM Generation** → stream/oneshot via the broker
9. **Confidence** → score vs threshold (PASS/repair)

Quick mode (default) skips 2–4/HyDE and uses the AgentBus directly; the orchestrator is the "deep" path.

---

8. Inference layer (`eli/cognition`)

- **`inference_broker.py`** — the **only canonical inference path**. `broker.infer()`.
- **`gguf_inference.py`** — llama-cpp wrapper, live runtime params, `chat_completion()`, runtime snapshot publishing. Model-agnostic; loads whatever GGUF is configured.
- **`reasoning_modes.py`** — modes: `quick`, `chain_of_thought`, `self_consistency`, `tree_of_thoughts`, `constitutional_ai` (private execution strategy, not shown raw).
- **`chat_model.py`, `context_builder.py`, `context_synthesiser.py`, `user_info_builder.py`** — prompt/context assembly.
- **`llm_intent.py`** — LLM fallback for ambiguous routing. Decoding is constrained by a GBNF grammar built from the live action catalogue, so the model can only name a capability that actually exists and can only emit well-formed JSON — an invented action is unrepresentable rather than rejected after the fact. Backends without grammar support (Ollama, remote) fall back to the free-text path unchanged.
- Token/context budgeting lives in `engine.py::_build_enhanced_system` (budget-aware persona trim) + `core/dynamic_runtime_budget.py`.

---

9. Grounding & anti-confabulation spine (`eli/runtime` + `eli/core/grounding.py`)

ELI's signature subsystem — keeps a local model from confidently stating unbacked facts.

- **`deterministic_grounding_gate.py`** (4.3k) — the gate; decides when an answer must be backed by evidence vs may be free CHAT; blocks raw evidence/telemetry packets from leaking to the user.
- **`grounding_escalation.py`** — *tiered* escalation (built this cycle): a checkable factual turn with **low grounding** escalates — external fact → web agent, self/project fact → broad local agent fan-out — and **hedges honestly** if nothing grounds it. Trigger is grounding, *not* the (often-high) response-confidence. Env: `ELI_GROUNDING_ESCALATION`.
- **`evidence_ledger.py`, `evidence_store.py`, `evidence_arbitration.py`, `memory_evidence.py`** — what counts as evidence, where it's recorded, conflict resolution.
- **`persistence_gate.py`** — stops internal report-dumps/junk being written to memory (and replayed later).
- **`response_contracts.py`, `response_packets.py`, `response_policy.py`, `final_response_assembly.py`, `final_response_provider.py`, `user_visible_response_surface.py`** — shape/clean the final user-visible answer; format internal "surface packets" into readable text (never raw JSON).
- **`truth_report.py`, `eli_identity_audit.py`, `control_contracts.py`** — runtime truth evidence, identity grounding, control-action contracts.
- **`output_governor.py`, `response_governance.py`, `response_sanitizer.py`, `persona_hygiene.py`** (cognition) — final governance/sanitisation pass.

This is also where the **known weak seam** lives: internal-state→spoken-output leaks (identity confabulation, telemetry dumps) — partially defended, the active frontier (see §20).

---

10. Execution layer (`eli/execution/executor_enhanced.py`)

- `execute(action, args) -> dict` — **216 dispatch actions**; runtime capability manifest reports **216 capabilities, 199 of them routable** *(live, read from the manifest each boot)*.
- **PHASE45 fast-path** (`engine.py`) — deterministic OS/media/status/job actions (`VOLUME`, `MEDIA_CONTROL`, `NEXT_MEDIA`, `OPEN_APP`, `DATE`, `SHELL_EXEC`, `ANALYZE_IMAGE`, `CHECK_JOB`, `BACKGROUND_JOBS`, …) bypass the LLM and return the executor result **verbatim** — never paraphrased.
- Action families: media/OS control, file/document (`SUMMARIZE_FILE`, `ANALYZE_PDF[_FOLDER]` incl. per-file mode, `CREATE_DOCUMENT`), web (`WEB_SEARCH`), news, weather, memory (`MEMORY_STORE/RECALL`), code (`CODE_SOLVE`, `GENERATE_SCRIPT` → §coding), background jobs, self/runtime reports (`RUNTIME_STATUS`, `SELF_REPORT`, `RUNTIME_AUDIT`, `EXPLAIN_*`), vision, image generation, scheduling/reminders.
- **Background jobs** (`runtime/background_tasks.py`) — `submit/get/list`; heavy PDF-folder analysis auto-backgrounds; `CHECK_JOB` surfaces the real result.

Plugins (`eli/plugins`, manager.py)

Bundled: `calendar`, `document_reader`, `media`, `notes`, `pomodoro`, `system_stats`, `tts`, `weather`, `web`, `web_automation`. State in `config/plugins_state.json` (gitignored). Custom agents/plugins are trust-gated.

---

11. Memory subsystem (`eli/memory`)

- **`memory.py`** (`Memory`, `get_memory()`) — long-term store over SQLite + **FTS5**, plus failure/observation/habit logging with anti-junk filters.
- **`vector_store.py`** — **FAISS** index at `artifacts/vectors/index.faiss`; embedder `models/embeddings/nomic-embed-text-v1.5.Q4_K_M.gguf` *(runtime)*.
- **Knowledge graph** — `kg_entities` / `kg_relations` (FTS5-backed) in the user DB.
- **Working memory** (`cognition/working_memory.py`) — session-pinned facts, junk-pin guard; restored across sessions.
- **Persona/personal memory** — `runtime/personal_memory_*`, `persona_updater`.

Two databases *(runtime)*: `artifacts/db/user.sqlite3` (user-facing) and `artifacts/db/agent.sqlite3` (agent/self-improvement). Observed tables include: `memories(+_fts)`, `conversation_turns`, `conversations`, `kg_entities(+_fts)`, `kg_relations`, `recall_log`, `runtime_events`, `news_articles(+_fts)`, `news_reflections`, `habits/habit_events/habit_rules`, `observations`, `learning_replay`, `working_memory_pins`, `user_patterns`, `session_summaries`, `corrections`, `failures`, `error_tracking`, `improvements`, `capability_proposals`, `code_patches`, `agent_dispatches`, `agent_metrics`. Retrieval is hybrid: keyword + FTS5 + FAISS + RAG + KG, merged & reranked (§7).

---

12. Perception (`eli/perception`)

- **Vision** — `vision.py` (model-agnostic VL: Moondream / Qwen2.5-VL hot-swap; mtmd encoder forced to CPU to avoid CUDA clip segfault), `analyze_image.py`, `ambient_vision.py` (ambient toggle), `screen_locator.py`, `gaze_engine.py`, `analyze_mesh.py`, `extract_equations.py`.
- **STT** — `local_whisper_stt.py` (faster-whisper `small.en`, CPU int8, `local_only` when offline), `audio_stt.py`, `eli_listen.py`, `voice_worker*.py` (wake-word “computer”, music-bleed filter, echo gate).
- **TTS** — `tts_router.py` (Piper voices, e.g. `en_US-amy-medium`).
- **OS** — `os_controller.py` (app/window/keyboard/mouse), `analyze_csv.py`, `analyze_pdfs.py`.

---

13. Persona system (`eli/cognition/persona*`)

- `persona.txt` (authored) + `persona.auto.txt` (overlay) — the voice.
- `persona_updater.py` — updates overlay/KG/user-profile each turn (tone signals).
- `persona.py`, `persona_values.py`, `persona_status.py`, `persona_hygiene.py`.
- `tone_analyzer.py` — writes tone preference signals (correction/humor/depth).
- Generation injects a budget-trimmed persona + situation brief (`engine.py::_build_enhanced_system`); a live-runtime fact is injected for model/identity questions so ELI reports the loaded model instead of denying it.

---

14. Daemons & background loops

- **Proactive daemon** (`planning/proactive_daemon.py`) — continuous learning; pattern signals (time habit, topic focus, recurring errors, active project).
- **Self-improvement loop** (`runtime/self_improvement.py`) — learns from logged failures (now filtered of transient/user errors).
- **Habit scheduler** (`planning/habits_scheduler.py`, `habits.py`) — learned routines; guarded against junk-rule replay.
- **Scheduler / task bus** (`kernel/scheduler.py`, `kernel/task_bus.py`).
- **Background tasks** (`runtime/background_tasks.py`) — async heavy jobs.
- **Code monitor** (`runtime/code_monitor.py`), **ambient vision loop**.
- **World event bus** (`world/world_event_bus.py`) — fed confidence/agent events.

---

15. Learning / self-training (`eli/learning`)

LoRA fine-tuning pipeline on a Phi-3 base: `bootstrap_phi3_base.py`, `base_model_resolver.py`, `dataset_builder.py`, `dataset_filters.py`, `merge_reviewed_datasets.py`, `export_trainable_dataset.py`, `training_preflight.py`, `lora_trainer.py`, `lora_trainer_guard.py`, `lora_eval.py`. Turns reviewed interactions into trainable datasets and adapters.

---

16. EliWorld (`eli/world`, `eli/gui` EliWorld tab, `kernel/world_model.py`)

An internal/spatial "world" the agent inhabits (rooms like the Reflection Chamber/Anomaly Room appear *(runtime)*). `world_event_bus.py` receives runtime events; `local_world_bridge.py` bridges to the model; snapshots/ledger/journal under `artifacts/world/`. Experimental/creative subsystem.

---

17. Core infrastructure (`eli/core`)

- **`netguard.py`** — offline-by-default: `guarded_urlopen`, process-wide socket guard (fail-closed on non-loopback while offline), `allow_network()` scoped context for deliberate user-initiated fetches (e.g. model download).
- **`paths.py`, `portable_paths.py`, `legacy_paths.py`, `db_paths.py`** — canonical path resolution (dev vs platformdirs; `ELI_*` overrides).
- **`runtime_settings.py`, `config.py`** — settings load/merge/heal; `DEFAULTS` are clean (offline, no model, wizard on). `config/settings.json` is per-user/gitignored, seeded from `config/templates/settings.template.json`.
- **`hardware_profile.py`, `startup_hardware_optimizer.py`, `dynamic_runtime_budget.py`** — detect GPU/VRAM, pick ctx/gpu_layers/batch.
- **`first_run.py`, `first_run_wizard.py`** — onboarding state.

- **`model_download.py`** — curated GGUF downloader (catalog, resumable, GGUF-magic + size validated, netguard-gated). Install-time menu only — inference stays model-agnostic.
- **`grounding.py`** — `is_grounded_query` etc. (shared grounding helpers).

---

18. GUI (`eli/gui`)

PySide6 (LGPL — not PyQt/GPL, see `pyproject.toml`). `eli_pro_audio_gui_v2_0.py` (main window, 11k LOC), `app.py` (entry/boot), `labs_tab.py` (Experimental), panels in `eli/gui/panels/` (`startup.py` = StartupModelSelectionDialog + FirstBootWizard with curated download, HardwareTuningDock). EliWorld tab, Operator console, Proactive dock. `qt_compat.py` shim allows PyQt fallback from source.

---

19. On-disk layout (`artifacts/`, runtime)

```
artifacts/
├── db/{user,agent}.sqlite3        # memory + agent state
├── vectors/index.faiss            # semantic index
├── conversations/ , conversations/archive/
├── incidents/<date>.jsonl         # incident log
├── proactive/ , world/{snapshots,ledger,journal}/
├── runtime/users/<uuid>/user_profile…   # per-user profile
├── runtime_snapshot.json          # live model/runtime truth
├── documents/ , scripts/          # generated artifacts
└── analyze_image_*/               # vision outputs
config/  settings.json(gitignored) · templates/ · plugins_state.json …
models/  <your>.gguf · embeddings/ · whisper/ · image/
voices/  Piper ONNX
```

---

20. Known seams & weak points (honest)

1. **God-files** (§1) — four files ≈ a third of the codebase; the active refactor target. Split by concern, hold a size ceiling, keep a repair-log.
2. **Internal-state → spoken-output seam** — the recurring failure class (identity confabulation, telemetry/packet leaks, ungrounded factual claims). The grounding gate + escalation defend it but don't fully close it on the plain CHAT path. This is *the* frontier item.
3. **Swallowed errors** — many `except Exception: pass`; convert to logged/raised so failures surface instead of hiding.
4. **No capability eval until now** — `tools/eval/` (see `tools/eval/README.md`) is the new measurement layer; coverage grows by turning every bug-log into a case.
5. **Latency vs model size** — on an 8GB GPU a 24B at Q5 offloads few layers → minutes/reply; a 7–9B Q4 is the sweet spot. The model picker should warn on poor VRAM fit.

---

21. Where to make a change (quick index)

To change…	Edit…
how input is routed to an action	`eli/execution/router_enhanced.py`
what an action does	`eli/execution/executor_enhanced.py`
which agents run / how grounded	`eli/cognition/agent_bus.py`
the deep 12-stage retrieval	`eli/cognition/orchestrator.py`

To change…	Edit…
prompt/persona budget, the engine spine	`eli/kernel/engine.py`
anti-confabulation / verify-or-hedge	`eli/runtime/grounding_escalation.py`, `deterministic_grounding_gate.py`
final answer shaping/cleanup	`eli/runtime/*response*`, `cognition/output_governor.py`
memory store/recall	`eli/memory/memory.py`, `vector_store.py`
model loading / inference	`eli/cognition/gguf_inference.py`, `inference_broker.py`
offline/network policy	`eli/core/netguard.py`
paths / settings / model download	`eli/core/{paths,runtime_settings,model_download}.py`
the GUI	`eli/gui/eli_pro_audio_gui_v2_0.py`, `gui/panels/`
eval / regression board	`tools/eval/` (+ `tools/eval/README.md`)

“‘