

ELI — Architecture Ascii

ELI v2.4.69

September 2026

Contents


Blueprint — ELI Full Architecture (ASCII)	1
A. Full system — layered block diagram	1
B. The 15 agents	3
C. Module tree (LOC · key files · role)	3
D. On-disk data layout	5
E. The flow in one line	5


Blueprint — ELI Full Architecture (ASCII)

Updated for v2.4.69. All CHAT modes → gradient orchestrator; bus composed at S06.

The entire system in one drawing, plus the module tree and data layout. Grounded in the real source (see `architecture.md` for prose, `diagrams.md` for the pipeline/memory/gating close-ups). Every layer and box maps to a real path.

A. Full system — layered block diagram





+ goal/scheduler ticks → proposals (observe-only / memory-write; need approval)
habits_scheduler/habits – routines scheduler – jobs
background_tasks – async heavy work reflection loop ambient_vision loop
scheduled_tasks (durable overnight/timed) world_event_bus ← confidence events

```

OUTPUT ← output_governor / sanitiser ← inference ← (any path)
          |
          └→ TTS (Piper) / GUI render

```

B. The 15 agents

AGENT BUS (cognition/agent_bus.py :: AgentBus.dispatch → DispatchResult)
selects a set → runs on a dependency DAG → aggregates grounding confidence

BusMemoryAgent	recall: SQLite/FTS5/FAISS → context
KnowledgeGraphAgent	entities/relations from the KG
SystemAgent	runs executor actions (WEB_SEARCH, RUNTIME_AUDIT, ...)
OrchestratorAgent	plan / coordinate grounded synthesis
FileCodeAgent	searches the codebase for code-grounded answers
IntrospectionBus...	live runtime + gathers identity/awareness audits as evidence (persona summarises – never a data dump)
CapabilityAgent	capability manifest (what ELI can do)
ReflectionAgent	session reflection / insights
ProactiveAgent	proactive patterns / signals
HabitAgent	learned habits
SelfImprovementAgent	self-improvement state
FrontierAgent	frontier reasoning hooks
PluginAgent	dispatch to registered plugins
VoiceAgent	STT/TTS engine state
CriticAgent	second-tier verifier: checks retriever agents' output
+ CodeAgent eli/coding/agent	SEPARATE pipeline: plan→search→verify→repair (CODE_SOLVE / GENERATE_SCRIPT) – not a bus agent

C. Module tree (LOC · key files · role)

```
eli/ (194,623 LOC, 437 files) · api/server.py (FastAPI web app + dashboard)
```

```

└─ __main__.py ..... entry dispatch (GUI | --headless)

└─ kernel/          17.5k ─ the core
    │ engine.py      16.0k  CognitiveEngine.process() – the spine
    │ scheduler.py    timed jobs
    │ pipeline.py · state.py · world_model.py · self_upgrade.py

└─ execution/       26.3k ─ route → act
    │ executor_enhanced.py 15.8k 205 supported / 227 manifest
    │ router_enhanced.py   8.3k  priority pipeline
    │ execution_planner.py · route_authority.py · route_contracts.py
    │ operator_actions.py · operator_policy.py

```

- └─ media_runtime.py
- └─ portable_intent_contract.py
- ─ cognition/ 20.8k — think
 - └─ agent_bus.py 3.5k 15 agents + dispatch
 - └─ orchestrator.py 12-stage deep retrieval
 - └─ gguf_inference.py 3.2k · inference_broker.py model-agnostic inference
 - └─ reasoning_modes.py · hyde.py · reranker.py · llm_intent.py
 - └─ persona.py/_updater/_values/_status/_hygiene (+ persona.txt, persona.auto.txt)
 - └─ context_synthesiser.py · working_memory.py
 - └─ output_governor.py · response_governance.py · response_sanitizer.py
 - └─ grounded_status.py · introspection_agent.py · tone_analyzer.py
- ─ runtime/ 35.4k — grounding spine + daemons (95 files)
 - └─ deterministic_grounding_gate.py 3.4k · grounding_escalation.py
 - └─ evidence_ledger/arbitration.py
 - └─ persistence_gate.py · truth_report.py · control_contracts.py
 - └─ response_contracts/packets/policy.py · final_response_assembly/provider.py
 - └─ user_visible_response_surface.py · personal_memory*.py · reflection.py
 - └─ background_tasks.py · self_improvement.py · code_monitor.py
 - └─ capability_sync.py · pending_proposal.py · runtime_policy.py
- ─ memory/ 8.8k — remember (11 files)
 - └─ memory.py 5.7k Memory · SQLite + FTS5 · storage policy in policy.py
 - └─ vector_store.py FAISS index
- ─ perception/ 10.2k — sense (23 files)
 - └─ vision.py · analyze_image/csv/pdfs/mesh.py · ambient_vision.py
 - └─ audio_stt.py · local_whisper_stt.py · voice_worker_streaming.py
 - └─ tts_router.py · os_controller.py · screen_locator.py · gaze_engine.py
- ─ planning/ 4.0k — proactivity (19 files)
 - └─ proactive_daemon.py · habits_scheduler.py · habits.py · jobqueue_cli.py
- ─ coding/ 2.1k — CodeAgent (plan→search→verify→repair, 12 files)
- ─ learning/ 4.3k — LoRA self-training (14 files)
 - └─ lora_trainer/eval/guard.py · dataset_builder/filters.py
 - └─ bootstrap_phi3_base.py · base_model_resolver.py · training_preflight.py
- ─ plugins/ 5.9k — manager + bundled plugins (34 files)
 - └─ calendar · document_reader · media · notes · pomodoro · weather · system_stats · tts · web · web_automation
- ─ world/ 1.8k — EliWorld (world_event_bus, local_world_bridge)
- ─ core/ 11.2k — infra
 - └─ netguard.py offline failsafe + allow_network()
 - └─ paths.py · portable_paths.py · legacy_paths.py · db_paths.py
 - └─ runtime_settings.py · config.py · grounding.py
 - └─ hardware_profile.py · startup_hardware_optimizer.py · dynamic_runtime_budget.py
 - └─ model_download.py
- ─ gui/ 27.4k — PySide6 desktop
 - └─ eli_pro_audio_gui_v2_0.py 13.1k · app.py · labs_tab.py 5.7k

```

├── panels/ (startup.py: model picker + FirstBootWizard, HardwareTuningDock)
├── tools/ 7.6k — image_engine · news · document tools
├── contracts/ 0.8k · cli/ 0.1k · system/ 1.2k · utils/ 1.8k · setup/ 1.9k · onboarding/ 0.7k

```

D. On-disk data layout

```

artifacts/
├── db/
│   ├── user.sqlite3 memories · memories_archive · memory_meta · semantic ·
│   │               conversation_turns · conversations · session_summaries ·
│   │               kg_entities · kg_relations · recall_log · runtime_events ·
│   │               learning_replay · observations · habits/habit_events/habit_rules ·
│   │               user_patterns · user_model · eli_stances · belief_revisions ·
│   │               news_articles · news_reflections (+ FTS5 indexes)
│   ├── agent.sqlite3 agent_dispatches · agent_metrics · improvements · failures ·
│   │               corrections · capability_proposals · error_tracking
│   ├── system_index.sqlite3 OS index
│   └── coding_memory.sqlite3 coding_bug_fixes
├── vectors/index.faiss semantic index (embedder: nomic-embed-...Q4_K_M.gguf)
├── conversations/
├── runtime/ · runtime_snapshot.json live model/runtime truth
├── world/{snapshots,ledger,journal}/
├── documents/ · scripts/ · analyze_image_*/
├── config/ settings.json (gitignored) · templates/settings.template.json ·
│   settings.example.json · plugins_state.json
├── models/ <your>.gguf · embeddings/ · whisper/ · image/ tts_piper/ Piper voices

```

E. The flow in one line

```

INPUT → ROUTER → ENGINE └─ FAST-PATH ────────────────────> OUTPUT (verbatim)
                        └─ ORCHESTRATOR (all CHAT modes, gradient)
                            ├── retrieve_for_turn() → MEMORY
                            ├── dispatch_specialists() → 15-agent BUS
                            └─ broker.infer() → finalize_turn() (S12)
GATES: netguard · persistence · grounding · escalation · governor
INFERENCE (broker → gguf, model-agnostic) → OUTPUT → TTS/GUI
BACKGROUND: proactive · self-improve · habits · learning(LoRA) · world · scheduler

Companion docs: architecture.md (prose, every subsystem) · diagrams.md (pipeline / memory
/ gating close-ups) · tools/eval/README.md (the measurement layer). “”

```