

ELI — Common Errors and Fixes

ELI v2.5.16

October 2026

Contents

ELI — Common Errors & Fixes **2**
ELI stopped starting after a system upgrade: `No module named ...` 3
"ELI's window cannot start: This system has no libGL.so.1" 3
`undefined symbol: g_variant_builder_init_static` / `Failed to load module: ...libgvfsdbus.so` 3
The first-run questions took a command as my answer ("Your name is resume") 3
macOS: `install.sh: line ...: _PIP_LINKS[@]: unbound variable` 4
Windows: the installer stops at "Upgrading pip", or ends with `UnicodeEncodeError: 'charmap'` 4
The Linux AppImage does nothing, or: `version 'GLIBC_2.38' not found` 4
From-source install on a clean system: stops early, or "complete" with no GUI 4
Voice is silent although a `piper` program is installed 5
GPU pack: `libcudart.so.12: cannot open shared object file` 5
GPU pack: `no CUDA wheel found` while offline-by-default 5
Windows: two ELI windows after Setup.exe 5
Setup wizard: nomic / Piper "won't download" 6
Web server / phone "Home" page won't connect 6
Microphone / voice input silent or "ELI can't hear me" (desktop) 6
YouTube won't play / "mpv is not installed" / no video window 6
Spotify: "play " opened the search but played nothing 7
AppImage: an app ELI opened closes at once, or Spotify/yt-dlp calls fail 7
Startup: "could not verify your settings in time" 7
ELI quits with "Fatal Python error: none_dealloc: deallocating None" 7
Wayland: mouse clicks do nothing / ydotool errors 8
Clicking a link opens no browser page 8
Wrong / tiny model loads instead of your usual one 8
Server won't start / "address already in use" 8
Phone can't connect after a restart 9
Phone can't connect at all (fresh pairing) 9
Phone microphone / voice doesn't work 9
Model dropdown in the dashboard is empty 9
News shows old stories as "the latest" 9
Terminal shows escape-code junk around a URL (`^[]8;;…`) 9
`git push` → "Could not resolve host github.com" 9
Vision / CLIP segfaults on the GPU 9
Arch / lean distros: "attempt to write a readonly database" (portable install) 10
Arch / lean distros: GUI won't open — "libxcb-cursor0 is needed to load the Qt xcb platform plugin" 10

Running ELI on Arch — verified working steps . . . 10
“ELI can’t see my Ollama models” (any OS) . . . 11
ELI dies mid-session: `double free or corruption` / `Aborted (core dumped)` . . . 11
The terminal floods with `Audio source must be entered before listening` . . . 12
ELI reports facts about you, then calls them a hallucination when you push back . . . 12
A reply ends abruptly on a dangling - . . . 12
“ELI refuses a general question with ‘the net’s off, I can’t verify that’ ” . . . 12
Natural voice / cloning: `pip install "eli-v2.0[natural]"` fails with “No matching distribution” 13
Natural voice / cloning: `ImportError: cannot import name 'isin_mps_friendly'` . . . 13
Natural voice / cloning: hangs or `EOFError` on first use . . . 13
Natural voice / cloning: `ModuleNotFoundError: No module named 'torchaudio'` (or `torchcodec`) . 13
A community plugin won’t install . . . 14
A plugin is installed but does nothing . . . 14
A plugin asked for nothing and was denied anyway . . . 14
An MCP server was added but has no tools . . . 14
“ELI is offline” but a plugin still reached the internet . . . 15
Training says my GPU can’t do it . . . 15
“Install” opened a review dialog instead of just installing . . . 15
A plugin is enabled but raises PermissionError . . . 15
A registry or download was “refused for safety” . . . 15
ELI did something I did not ask for, in the middle of a sentence about something else . . . 15
ELI said it had done something and it had not . . . 16
“open nasa website” answered “not installed, install it?”; “open that website” searched for the words . . . 16
“maximise browser” changed ELI’s own window, or nothing came back . . . 16
I gave ELI a folder and it asked me to run `ls` and paste the output . . . 16
A short remark took half a minute before the reply started . . . 17
Normal mode takes two or three minutes to answer even a short remark . . . 17
Voice input never works: `transcription failed: TypeError: open() got an unexpected keyword argument 'metadata_errors'` . . . 17
The first message after start-up takes minutes and the model loads a second time . . . 17
ELI’s voice: `PortAudioError: Error opening OutputStream: Invalid sample rate` . . . 17
ELI started `ollama serve` at start-up while a GGUF model was in use . . . 18
A library reported usage to its maker . . . 18
Asked why a reply was slow, ELI gave RAM figures . . . 18
“go deeper” on a story from the news briefing got “that story doesn’t exist” . . . 18
“Nothing was carried out in this turn.” under a reply to a greeting . . . 19
A question about ELI got an invented explanation . . . 19
The next message waited behind background work . . . 19
The runtime report said the context window was “reduced” to a larger number . . . 19
v2 used folders and files other ELI versions use too . . . 19

ELI — Common Errors & Fixes

> **Updated for v2.4.84 (September 2026).** Primary path: GitHub Releases (`ELI-Setup-*.exe`, macOS dmg, Linux AppImage, portable `./ELI_Setup.sh`). Wizard core stages hard-fail until nomic + voice + chat model are present. Windows: one desktop ELI + GUI singleton. **Home** tab ≠ Home Assistant.

A quick reference for the gremlins that actually come up when running ELI. Each one is symptom → cause → the exact fix. Add new ones as they surface.

ELI stopped starting after a system upgrade: `No module named ...`

Symptom (from-source install, any OS): ELI worked, the operating system or Python was upgraded, and now every start ends in `ModuleNotFoundError: No module named 'requests'` (or `eli`, `PySide6`, ...). On Windows the message is `No Python at '...\python.exe'`. **Cause:** ELI's workspace (`.venv`) is tied to the Python it was built with. An upgrade that replaces that Python (Ubuntu 24.04 to 26.04 swaps 3.12 for 3.14; uninstalling or upgrading Python on Windows or macOS does the same) leaves the workspace starting as the new version, which cannot see a single installed package. **Fix:** start ELI the usual way. The launcher now says what happened and, if the old Python version is still installed anywhere (PATH, the `py` launcher, uv, pyenv, Homebrew), points the workspace back at it with nothing reinstalled. If it is not, run the installer once:

```
bash install.sh          # Linux / macOS
install.bat              # Windows
```

It rebuilds the workspace on the Python you have now. Models, memory and settings are not touched. To check or repair by hand: `python3 scripts/eli_env.py status` / `repair`. The release packages (AppImage, Setup.exe, dmg) bring their own Python and are not affected.

"ELI's window cannot start: This system has no libGL.so.1"

Symptom: on a server, a container or a very minimal install the window does not open. Older versions printed `Please install PySide6`, which was not the reason. **Cause:** the graphics libraries (`libGL`, `libEGL`) belong to the GPU driver and the display server. ELI cannot carry its own copy without breaking hardware acceleration, so it uses the system's, which every desktop has. **Fix:** run the command in the message, for example `sudo apt-get install -y libgl1 libegl1`. `install.sh` adds these itself when it can and the `.deb` depends on them. The server (`ELI-Server`) and terminal modes do not need them.

`undefined symbol: g_variant_builder_init_static` / `Failed to load module: ...libgvfsdbus.so`

Symptom: warnings in the terminal at every start of the AppImage on a distribution newer than the one it was built on (Ubuntu 26.04, Fedora, Arch); file dialogs without network places. **Cause:** the AppImage carried GTK and GLib's I/O library from its build machine, brought in by Qt's GTK theme plug-in. Those load their own plug-ins from the system (gvfs, dconf, input methods), which are built for the system's newer GLib and do not load into the older one. **Fix (this version):** the Linux bundle carries neither (`packaging/linux/trim_bundle.py` removes the two Qt plug-ins that bring them in and every library only they needed, and the release checks it). ELI draws its own widgets; a file dialog is the desktop's own, asked for through its portal, or Qt's where there is no portal. Making the AppImage use the system's GLib was tried earlier and withdrawn: GTK then failed to load on those same systems.

The first-run questions took a command as my answer ("Your name is resume")

Symptom: a new profile is asked what to call you; you type a command instead (`resume`, `louder`, `next`) and ELI stores it as your name. **Cause:** the interview decided what was a command from a short list of verbs of its own. **Fix (this version):** it asks the router, as every message is asked. A message the router turns into an action is carried out and is never an answer; a name introduced as one ("call me Pause") is a name whatever the word. To correct a stored name: `my name is ...`.

---

macOS: `install.sh: line ...: _PIP_LINKS[@]: unbound variable`

Symptom: `bash install.sh` stops at "Installing PyTorch" on a Mac; `scripts/eli_startup.sh` stops with `APP_ARGS[@]: unbound variable` when started with no arguments. **Cause:** macOS ships bash 3.2, which reads the expansion of an empty array as an unset variable under `set -u`. Every Linux has a newer bash, so it was never seen there. **Fix (this version):** every shell script expands arrays in the form bash 3.2 accepts, and uses nothing newer than it; the tests hold that. The installers themselves now run in CI on Linux, macOS and Windows (`installer-smoke.yml`).

Windows: the installer stops at "Upgrading pip", or ends with `UnicodeEncodeError: 'charmap'`

Symptoms: `ERROR: To modify pip, please run the following command: ... python.exe -m pip install --upgrade pip`, then `pip failed`; or the installation completes and the last step prints `UnicodeEncodeError: 'charmap' codec can't encode character '\u2192'`. **Cause:** `pip.exe` cannot replace itself while it runs, and the installer called it that way. And any output that is piped or captured on Windows uses the legacy code page, which has no arrow or emoji: the first one printed raised. **Fix (this version):** pip is always run as `python -m pip`; importing ELI makes its output streams write `?` for what they cannot show (the packaged apps already did); the installer reports its own result, not its last optional step's. It also installs within the locked versions from the start, so no newer Qt add-on packages are left beside the locked Qt, and it falls back to version ranges, then to one requirement at a time, as `install.sh` does.

The Linux AppImage does nothing, or: `version 'GLIBC_2.38' not found`

Symptom: double-clicking the AppImage does nothing. From a terminal: `Failed to load Python shared library ... libm.so.6: version 'GLIBC_2.38' not found`. **Cause:** a Linux program only starts where the system library (glibc) is at least as new as the one it was built against. The AppImage is built on Ubuntu 24.04 and needs **glibc 2.39** or newer: Ubuntu 24.04+, Debian 13+, Fedora 40+, Linux Mint 22+, Pop!_OS 24.04+, RHEL / AlmaLinux / Rocky 10+, openSUSE Leap 16+ and Tumbleweed, Arch. Ubuntu 22.04 (2.35), Debian 12 (2.36), Mint 21 and RHEL 9 (2.34) are older. Check with `ldd --version`. **Fix:** on an older system install from source, which builds against the libraries the system has: download `ELI_v2-<v>-linux-portable.tar.gz` from the same release, unpack it and run `./ELI_Setup.sh`. Nothing can be added to the AppImage to make it start; upgrading the distribution also works. Current AppImages say this themselves, in a sentence and a dialog, instead of the loader's message.

From-source install on a clean system: stops early, or "complete" with no GUI

Symptoms (a clone or the source package, Linux): - `The virtual environment was not created successfully because ensurepip is not available` at the first step (stock Debian and Ubuntu). - `bash install.sh` ends without a message right after `Installing ELI (editable)`, exit status 2 (any clone: there is no built wheel in `dist/`). - `ERROR: Failed building wheel for PyAudio` (`Python.h: No such file or directory`), the installer prints "installation complete", and ELI then says `Please install PySide6`.

Cause: three separate faults on a system that has only Python. Debian and Ubuntu package the part of Python that makes environments separately. A lookup for a prebuilt wheel ended the script when there was none. And pip installs a set of requirements all or nothing: PyAudio has no ready-made Linux build and needs Python's headers and PortAudio's to compile, so without them none of ELI's dependencies were installed, while the final check passed because it imported ELI from the folder it was run in. **Fix (this version):** the installer first works out what it builds with and is missing, adds it when that takes no password and otherwise prints one command, for example

```
sudo apt-get install -y python3.12-venv python3.12-dev build-essential portaudio19-dev
```

then installs ELI itself whatever happens to its dependencies, falls back from the pinned set to version ranges and then to one requirement at a time (naming any that failed), and checks the result from outside the source folder. An install that ended with errors now says so and exits non-zero. On an older checkout: run the command above, `mkdir -p dist`, and run the installer again.

Voice is silent although a `piper` program is installed

Symptom: no speech, and the log shows `piper failed rc=1` or a traceback from a `piper` outside ELI's folder. **Cause:** more than one program is called `piper`. Some Linux distributions ship an unrelated one that configures gaming mice, and a `piper` installed for a Python that has since been removed still sits on PATH. ELI used to take the first one it found by name. **Fix (this version):** ELI uses the Piper that was installed with it, and tries any other only if it starts and is the speech engine. Nothing to do. `ELI_PIPER_BINARY=/path/to/piper` still forces a particular one.

GPU pack: `libcudart.so.12: cannot open shared object file`

Symptom: AppImage or portable dies (or `eli` CLI dies) with `Failed to load … libllama.so: libcudart.so.12: … No such file or directory`. Disk shows `~/.local/share/ELI_v2/runtime/gpu` ~2.7 GB with `libggml-cuda.so` but **no** `libcudart.so.12` / `libcublas.so.12`. Iris Xe / AMD may still log `Vulkan backend forced` while installing a CUDA wheel. **Cause (≤2.4.25):** the `--vulkan` / Intel / AMD picker preferred `cuda-` gpu-packs assets, labelled them `vulkan`, and skipped NVIDIA runtime vendoring. **Fix (v2.4.26+):** upgrade; remove the broken pack and relaunch so the correct backend is downloaded:

```
rm -rf ~/.local/share/ELI_v2/runtime/gpu
# Launch the new AppImage (or: AppImage --integrate  then  eli)
```

NVIDIA machines get a CUDA pack **with** vendored cudart/cublas (or install fails closed). Intel/AMD/`--vulkan` get a real `vulkan-` wheel only.

Also check `which eli` — an old portable wrapper can still be first on `PATH` even after installing a newer AppImage. Run `./ELI_v2-*-x86_64.AppImage --integrate` so `~/.local/bin/eli` points at the AppImage you just installed.

GPU pack: `no CUDA wheel found` while offline-by-default

Symptom: Startup dialog: *GPU pack install failed: no CUDA wheel found…* Logs show `index cu124 unavailable (network disabled (offline mode): blocked connection to ('abetlen.github.io', 443))` on a machine that *does* have NVIDIA. **Cause (≤2.4.25):** GPU pack install used raw `urlopen` without a scoped `allow_network()` window, so NetGuard blocked the wheel index. **Fix (v2.4.26+):** install opens `allow_network("gpu-pack install")` and falls back to CI `cuda-` packs on the `gpu-packs` GitHub release if abetlen is empty. On portable, if `ggml_cuda_init` already lists your GPU, inference can still use the venv CUDA build — the dialog failure is about the optional GPU *pack*, not that CUDA is impossible.

Windows: two ELI windows after Setup.exe

Symptom: Finish the installer and two GUI windows appear (or desktop has ELI + ELI Server). **Cause:** Desktop used to pin both `ELI.exe` and `ELI-Server.exe`; post-install could race. **Fix (v2.4.23+):** Desktop shortcut is **ELI only**; Server stays in Start Menu. Frozen GUI takes a singleton lock — a second launch shows "ELI is already running" and exits. Close the extra window; upgrade to 2.4.23+ (current **2.4.84**).

Setup wizard: nomic / Piper "won't download"

Symptom: Chat model downloads, but embedder/voice stay missing. **Cause:** Older wizard called `download_aux(required_only=False)` and ignored failures; offline-by-default blocked FirstBoot prefetch. **Fix (v2.4.23+):** Wizard uses `python -m eli.core.model_download --aux` with a real exit code. Click **Fetch embedder + voice now** (scoped network) or run:

```
.venv/bin/python -m eli.core.model_download --aux
.venv/bin/python -m eli.runtime.voice_assets
```

Web server / phone "Home" page won't connect

Symptom: Settings → Web Server starts, phone can't open the URL / QR is `<this-computer-ip>`. **Cause:** LAN IP detection lacked a Windows path; Apply-port left a stale `ELI_API_PORT`; HTTPS sidecar couldn't stop; Windows firewall tip omitted 8443. **Fix (v2.4.23+):** Shared `resolve_lan_ip()` (UDP + PowerShell); Stop clears port env; HTTPS handle is retained; firewall hints include HTTPS. Use **phone / Wi-Fi** mode, allow the firewall command, open the **token** URL (Home tab = devices/MQTT, not Home Assistant).

Microphone / voice input silent or "ELI can't hear me" (desktop)

Symptom: ELI starts, `[AUDIO] Listening...` appears, but nothing is transcribed — or you had to set `ELI_MIC_DEVICE_INDEX` manually on v2.3.30. **Cause (v2.3.30 regression):** on Linux/PipeWire, PortAudio's `pulse` wrapper often passed a short probe yet stayed silent during real STT; on all platforms a webcam or virtual "Stereo Mix" default could win over a connected USB/Bluetooth headset. **Fix (v2.3.92+):** upgrade to **v2.3.92** — cross-platform ranked auto-resolve probes USB, headset/Chat, and Bluetooth endpoints before built-in/webcam/virtual routes, at each device's native sample rate. No env vars required on Linux, macOS, or Windows. **Diagnostics:**

```
python -m eli.tools.mic_diag          # lists devices + measures speech vs ambient
```

Manual override (only if auto-resolve picks the wrong endpoint):

```
python -m eli.tools.mic_diag   # note the [N] of your working device
export ELI_MIC_DEVICE_INDEX=N  # use YOUR index from mic_diag — not a fixed number
# Linux PipeWire/Pulse pin (optional):
export ELI_MIC_PULSE_SOURCE="$(pactl get-default-source)"
```

On Windows gaming headsets, set the **Chat** mic (not Game) as the default recording device in Sound settings. On macOS, pick your AirPods or headset under System Settings → Sound → Input.

YouTube won't play / "mpv is not installed" / no video window

Symptom: "play X on YouTube" fails, opens browser only, or plays audio with no video. **Cause:** `mpv` and/or `yt-dlp` missing. `yt-dlp` is a Python package (installed into ELI's bundled interpreter); `mpv` is an OS package. **Fix:** Say **yes** when ELI offers to install after a failed attempt, or run manually:

```
# Linux
sudo apt install mpv   # or dnf/pacman equivalent
# yt-dlp is installed by install.sh into ELI's Python — or: python -m pip install yt-dlp
# macOS
brew install mpv
# Windows
winget install mpv.MPV
```

Note: Plain “play X on youtube” uses **visible mpv**. “play X on youtube.com” opens the browser.

Spotify: “play ” opened the search but played nothing

Symptom: on a Wayland desktop (GNOME 50 has no X11 session) every “play X” took about 16 s and ended with “I opened the Spotify search … couldn’t confirm it started playing”. **Cause:** the only way ELI had to start a named song was typing into Spotify’s search box. Spotify runs there as a native Wayland window: xdotool/wmctrl cannot see it, so nothing could focus it or confirm focus, and ELI (rightly) refused to type blind. **Fix (in the code):** ELI asks Spotify over MPRIS to open `spotify:search:<song>`, which makes the desktop client play its top result with no window focus, then checks the song that started is the one asked for. If Spotify’s top result is a different song, ELI says what is playing. Typing into the search box is now only tried on X11 or for an XWayland window.

AppImage: an app ELI opened closes at once, or Spotify/yt-dlp calls fail

Symptom: “open spotify” says “Opened app: Spotify” but nothing appears; “play X on spotify” names the wrong song every time; `dbus-send` fails with `LIBDBUS_PRIVATE_1.16.2 not found`, or the system `yt-dlp` fails on `libcrypto`. Seen from the AppImage; a source install works. **Cause:** a frozen build points `LD_LIBRARY_PATH` at its own libraries and sets `QT_PLUGIN_PATH`, `SSL_CERT_FILE` and similar to paths inside itself. Programs ELI started inherited that, so a system program loaded the bundle’s older libraries and died. **Fix (in the code):** every program from outside the bundle now gets the machine’s own environment (`LD_LIBRARY_PATH` as it was before ELI started, no paths into the bundle); ELI’s own helper processes keep theirs. “Opened app” is only said once the app is still running a moment later; otherwise ELI says it started and closed. Programs Qt starts itself (links opened in the browser) are not covered.

Startup: “could not verify your settings in time”

Symptom: at launch the log says `[GUI][LOAD] could not verify your settings in time (probe timed out after Ns ...)` and the model loads with fewer GPU layers than you set. **Cause:** your GPU layers are above what ELI measured as fitting the free VRAM, so before loading them it runs them once in a separate process with a prompt the size of a real conversation. That check did not finish within its time budget. It runs the prompt in chunks and the line says how far it got: “the model had not finished loading”, or “loaded in Ns; the test prompt was at N/M tokens after Ns (R tokens/s)”. Other programs keeping every core busy slow it most (measured: the layers left on the CPU ran about ten times slower), so a slow rate points at what else is running. **Fix:** free VRAM (another program holding the card is listed at startup), lower the GPU layers or context to the measured fit, give the check more time with `ELI_LOAD_PROBE_TIMEOUT=<seconds>`, or skip it with `ELI_LOAD_PROBE=0` to load your settings unchecked. A timeout is remembered for an hour, so the next launch within that hour does not wait again.

ELI quits with “Fatal Python error: none_dealloc: deallocating None”

Symptom: after a while (twenty minutes in one session) the window closes and the terminal shows `Fatal Python error: none_dealloc: deallocating None: bug likely caused by a refcount error in a C extension`, with the main thread in a Qt call such as `refresh_audit_tab`. **Cause:** PySide6 6.12.0 on Python 3.11 and earlier takes one reference to `None` away on every Qt call that returns nothing (`setText`, `setItem`, `addItem`, …). When the count reaches zero Python aborts. From Python 3.12 `None` cannot be freed, so source installs on 3.12+ never see it. **Fix (in the code):** ELI gives `None`, `True` and `False` a large reserve of references at startup on Python before 3.12 (`eli/core/singleton_refs.py`), so no extension can bring them to zero, and installs on 3.10/3.11 skip PySide6 6.12.0. The release packages bundle

Python 3.12 (they were on 3.11 up to 2.5.8), and the release self-test fails when the bundled Qt loses references.

Wayland: mouse clicks do nothing / ydotool errors

Symptom: MOUSE_CONTROL or screen locate clicks fail on GNOME/KDE Wayland; error mentions `ydotoold`. **Cause:** `xdotool` only reaches XWayland windows. Native Wayland needs `ydotool` + the `ydotoold` daemon with `/dev/uinput` access. **Fix:**

```
sudo apt install ydotool    # or your distro equivalent
ydotoold &                  # run daemon (may need uinput group)
```

Or use AT-SPI/OCR click paths for labelled controls. See `docs/CROSS_PLATFORM.md`.

Clicking a link opens no browser page

Symptom: the server prints the URLs fine and copy-paste works, but Ctrl/right-click → "Open Link" opens nothing. **Cause:** broken **snap Firefox** — a `mesa-2404` auto-refresh leaves Firefox's GPU content-mount stale (`/snap/firefox/…/gpu-2404-provider-wrapper` goes missing), so *launching* a fresh Firefox fails silently while an already-open window still works. **Fix:**

```
pkill -9 firefox
sudo snap disconnect firefox:gpu-2404
sudo snap connect firefox:gpu-2404 mesa-2404:gpu-2404
# if it's still stuck:
sudo snap refresh firefox
sudo snap refresh mesa-2404
```

Reopen Firefox once, then relaunch the server — clicks (and the auto-open) work again. *(Nothing wrong with ELI — the URLs and server are fine; it's the snap.)*

Wrong / tiny model loads instead of your usual one

Symptom: ELI boots on a small model (e.g. tinyllama-1B) even though settings point at your 35B A3B. **Cause:** a stray `ELI_GGUF_MODEL_PATH` / `ELI_MODEL_PATH` env var in your shell **overrides config** (it wins over everything). **Fix:**

```
echo $ELI_GGUF_MODEL_PATH        # if this prints a model path, that's the culprit
unset ELI_GGUF_MODEL_PATH ELI_MODEL_PATH
```

Or just switch live from the dashboard: **Settings → Model**.

Server won't start / "address already in use"

Cause: a previous instance is still bound to the port. **Fix:**

```
pkill -f "api/server.py"      # or find it: ss -ltnp | grep 8081
```

Then relaunch. ELI's web port is **8081**, HTTPS **8443**.

Phone can’t connect after a restart

Symptom: a paired phone gets 401 / “can’t access the server.” **Cause:** it’s carrying an old token (server restarted, or you rotated the token). **Fix:** the token now **persists** across restarts — just re-open the current link/QR from the **Connect** tab. To deliberately kick a lost phone off, hit **rotate** and re-pair.

Phone can’t connect at all (fresh pairing)

Cause: the firewall is blocking the port. **Fix:** run the exact `sudo ufw allow …` lines the server prints on startup, and make sure the phone is on the **same Wi-Fi**.

Phone microphone / voice doesn’t work

Cause: browsers block the mic (`getUserMedia`) on a plain `http://LAN-IP` page. **Fix:** start with `--https` and open the `https://…:8443/` link on the phone; accept the one-time self-signed “not private” warning.

Model dropdown in the dashboard is empty

Cause: the list endpoint was being shadowed by the OpenAI-compatible `/v1/models` route. **Fix:** already fixed — the list lives at `/v1/models/installed`. If a fork regresses it, make sure the dashboard fetches that path, not `/v1/models`.

News shows old stories as “the latest”

Cause: the interest-matched half wasn’t recency-gated. **Fix:** already fixed (a freshness gate drops stale niche matches). If it recurs, say `refresh the news`.

Terminal shows escape-code junk around a URL (`^[]8;;…`)

Cause: you’re looking at a **log file** or `cat -v` output — those *always* render raw escape codes. A live terminal renders the link normally. Not a bug.

`git push` → “Could not resolve host github.com”

Cause: a transient DNS/network blip (ELI is offline-by-default, but git is separate). **Fix:** just retry the push.

Vision / CLIP segfaults on the GPU

Cause: the CLIP/vision path can segfault on some GPUs. **Fix:** CLIP runs on **CPU** by design; keep it there. Main-model + vision hot-swap handle VRAM automatically — don’t force CLIP onto the GPU.

Arch / lean distros: “attempt to write a readonly database” (portable install)

Symptom: on the portable Linux build, first-run DB init (Step 5) fails for the `user.*` stores with `sqlite3.OperationalError: attempt to write a readonly database`, while `system_index` / `coding_memory` / `agent` succeed in the same folder; the GUI then can’t open the memory DB and exits. **Cause:** the folder you extracted ELI into lives on a filesystem whose locking SQLite’s **WAL journal** needs isn’t supported — **NTFS, exFAT/FAT, or a network mount** (very common for a `~/Downloads` on a dual-boot data drive). WAL returns `SQLITE_READONLY` there; the rollback-journal stores beside it work fine, which is the tell-tale split. (This is also why the AppImage — which stores data under `~/.local/share/ELI_v2` on your main partition — worked while the portable build in `~/Downloads` didn’t.) **Fix:** fixed in **v2.1.19** — ELI detects a WAL-hostile filesystem and falls back to a DELETE (rollback) journal automatically, so it runs wherever you put it. On an older build, extract ELI onto an **ext4/btrfs** path (e.g. under your home partition) instead. Confirm your mount with: `findmnt -T . -o TARGET,FSTYPE`.

Arch / lean distros: GUI won’t open — “libxcb-cursor0 is needed to load the Qt xcb platform plugin”

Symptom: the AppImage’s engine loads (you see the `ELI v2.0` banner and the module log) but the window never appears; it aborts with the `xcb-cursor0 … xcb platform plugin` message. **Cause:** Qt 6.5+’s xcb platform plugin (`libqxcb.so`) links the **whole xcb-util family** — `libxcb-cursor`, `libxcb-icccm`, `libxcb-image`, `libxcb-keysyms`, `libxcb-render-util`, `libxcb-util` — which a minimal desktop doesn’t install. **Qt’s message is misleading:** it always blames `libxcb-cursor0` even when the real missing library is a different member (on a test Arch box the actual culprit was `libxcb-icccm.so.4`). **Fix:** fixed in **v2.1.21** — the AppImage now bundles the full xcb-util family (in `PySide6/Qt/lib`), so it launches out of the box on bare Arch/Debian with no extra packages. Since PySide6 6.11 the plugin also links `libxcb-shape` and the other xcb extension libraries; the build now reads the list from the plugin itself, bundles every one, and fails if one is missing. On an older build, install them yourself:

```
# Arch
sudo pacman -S xcb-util-cursor xcb-util-wm xcb-util-image xcb-util-keysyms xcb-util-renderutil xcb-util
# Debian/Ubuntu
sudo apt install libxcb-cursor0 libxcb-icccm4 libxcb-image0 libxcb-keysyms1 libxcb-render-util0 libxcb-util1
```

Diagnose exactly which library is missing (Qt’s own trace names it):

```
QT_DEBUG_PLUGINS=1 ./ELI_v2-*-x86_64.AppImage 2>&1 | grep -iE 'cannot|not found'
```

Running ELI on Arch — verified working steps

The **AppImage** is the easiest path: it bundles its own **Python 3.12**, so Arch’s system Python is irrelevant, and as of **v2.1.21** it bundles every Qt xcb library too. Download and run:

```
U=https://github.com/ShadowESC95/ELI_v2.0/releases/download/v2.4.84
wget "$U/ELI_v2-2.4.84-x86_64.AppImage"
chmod +x ELI_v2-2.4.84-x86_64.AppImage
./ELI_v2-2.4.84-x86_64.AppImage
```

Run it directly (as above) so it FUSE-mounts in place — `--appimage-extract-and-run` unpacks ~4 GB into `/tmp`, which fails with `libz.so.1: file too short` when `/tmp` is a small `tmpfs`. If you launch it over **SSH** into a running desktop session, prefix the command with `export DISPLAY=:0`. (Paste long URLs as a **single line** — a wrapped URL 404s and the tail runs as a stray command.)

Verified end-to-end on a clean Arch VM (XFCE): engine loads, database initialises on a WAL-hostile disk via the DELETE fallback, and the GUI opens with **every** xcb-util package uninstalled.

“ELI can’t see my Ollama models” (any OS)

Symptom: Backend is set to Ollama but the model list is empty, or ELI reports Ollama unreachable (`[Errno 111] Connection refused`) while Ollama is installed.

Fix (v2.1.30): ELI now starts Ollama for you. The #1 cause is simply that the local Ollama server isn’t running. If the `ollama` program is installed but stopped, ELI now runs `ollama serve` when you choose Ollama (as the backend, or an Ollama model in the toolbar or the start-up dialog), waits for it, and loads your models — no terminal needed. Since 2.5.15 refreshing the list alone never starts it, and with a GGUF model chosen ELI never starts or uses Ollama. If it still can’t start (Ollama not installed, or a remote host), you get per-OS guidance. Disable the auto-start with `ELI_OLLAMA_AUTOSTART=0`. Full walkthrough for non-technical users: `blueprints/ollama_startup_guide.md`.

Fix (v2.1.31): the toolbar dropdown itself was broken. Separately from the server ever starting, the toolbar’s Ollama dropdown could sit permanently empty with the status dot stuck on “Checking Ollama…” *even when Ollama was running and serving models perfectly*. The widget fetched the list on a background thread and handed the result back with `QTimer.singleShot(0, …)` — but a Qt timer started from a plain `threading.Thread` has no Qt event loop to run on, so the callback was silently dropped and the combo was never filled. It now hands results back over a Qt **signal**, which queues to the GUI thread correctly. The model-pull progress dialog had the same defect (progress never advanced, dialog never closed) and was fixed the same way. Regression-tested in `tests/test_ollama_selector_threading.py`, including a guard that fails if a future refactor reintroduces a timer on that path. Note this was invisible to the existing GUI tests because they only asserted the widget *constructs* — the new tests assert models actually land in the dropdown.

Earlier address-handling fixes (still in place) — you no longer need to do anything special about how the host is written:

- **`OLLAMA_HOST` set the way Ollama documents it** (`OLLAMA_HOST=127.0.0.1:11434`, with no `http://`) used to break ELI with `unknown url type`. Scheme-less hosts, bare IPs and missing ports are now all understood, everywhere — client, startup picker, wizard, Settings.
- **`localhost` resolving to IPv6 first** (common on Windows) while Ollama listens on IPv4 — ELI now automatically retries `127.0.0.1`, so this no longer looks like “Ollama is down”.
- **Ollama on another machine** was blocked by ELI’s offline-by-default guard with no explanation. A host you configure is now treated as a deliberate local service, like a LAN MQTT broker, so it is allowed and the Net toggle can stay OFF.

If it is still not found:

```
ollama list                    # is it running, and does it have models?
ollama pull llama3.2           # pull one if the list is empty
```

Per-OS start: **Windows/macOS** — Ollama starts itself (tray / menu bar); launch “Ollama” if it isn’t running. **Linux** — `systemctl --user start ollama`, or `ollama serve`.

For Ollama on a **different machine**, on that machine:

```
OLLAMA_HOST=0.0.0.0:11434 ollama serve   # listen beyond localhost
```

and allow port `11434` through its firewall. Then in ELI set **Settings → Model → Ollama → Host** to that machine (e.g. `192.168.1.20` — the port is filled in for you).

ELI dies mid-session: `double free or corruption` / `Aborted (core dumped)`

Symptom: the GUI is running fine, then the terminal prints a run of `Warning: maximum number 350 of (N_VOICES_LIST = 350 - 1) reached`, followed by `double free or corruption (!prev)` and `Aborted (core dumped)`. The whole app is gone.

Cause: not ELI. espeak-ng has a hardcoded `N_VOICES_LIST` of 350. On a machine with more installed voice/variant combinations than that, `espeak_ListVoices` walks off the end of its array and corrupts the heap. ELI reached it through `pyttsx3.init()` while listing system voices. Because that is a **native** `abort()`, the Python `try/except` around it caught nothing — the crash took the entire GUI with it.

Fix (v2.1.32): system-voice enumeration now runs in a **subprocess**. If espeak crashes, it kills a throwaway child; ELI logs the non-zero exit and continues with no system voices listed. Nothing to configure.

The terminal floods with `Audio source must be entered before listening`

Symptom: thousands of identical `[AUDIO] Listen error: Audio source must be entered before listening…` lines, a pinned CPU core, and a log that grows without limit.

Cause: `listen_once()` and the mic-calibration helper open `with self.microphone` on the **same** Microphone object the background listen loop is holding. `speech_recognition`'s `__exit__` sets `stream = None`, so when either finishes it closed the stream out from under the running loop — which then failed on every iteration, forever, at full speed.

Fix (v2.1.32): the loop detects the closed stream and **re-enters** it instead of spinning, and that error is rate-limited (once per 5s with a count) so no fault can bury the log again.

ELI reports facts about you, then calls them a hallucination when you push back

Symptom: ELI gives a grounded profile answer, you say "wrong", and it retracts everything — "that was likely a hallucination", "I don't have access to your personal files", "I only see the text in this window".

Cause: those statements are **false**. The facts came out of `user.sqlite3`. This is the base model's generic cloud-assistant reflex overriding what ELI actually is. There was a rule against *inventing* mechanisms, but none against *denying real ones* — the mirror failure.

Fix (v2.1.32): a `NO FALSE SELF-DENIAL` rule. Push-back is not proof ELI was wrong: it must correct the **specific** disputed field and say where the stored value came from, never retract a whole grounded answer as imaginary. The phrases also now trip the evidence validator. A companion rule stops "what do you know about yourself" being answered with *your* profile.

A reply ends abruptly on a dangling -

Symptom: a list-shaped answer stops mid-structure with an empty bullet on the last line.

Cause: the generator hit its token cap just as it opened the next bullet. The truncation detector missed it (a - is neither alphanumeric nor a sentence terminator, so it read as a clean ending), and the re-generation repair is deliberately **non-quick only** — a second inference would defeat quick mode's latency.

Fix (v2.1.32): the detector now recognises an orphan list marker, and the output governor trims it in **every** mode, so a capped answer ends on its last complete line.

"ELI refuses a general question with 'the net's off, I can't verify that' "

Symptom: an everyday advice/how-to question — *"what is the best way to eat breakfast"* — is met with *"I can't verify that right now — the net's off … turn the Net toggle on"*, and with the Net on it returns an irrelevant web snippet (a user saw *"Best Buy doesn't offer breakfast recommendations"*).

Cause: the fact classifier matched the "**what … is**" shape and treated the advice question as a checkable external fact (oddly, the apostrophe form *"what's the best way…"* slipped through and answered normally). So offline it hit the honest can't-verify hedge, and online it was web-searched — surfacing whatever the search returned.

Fix: two layers, no action needed on your part. - **v2.1.25** — advice / how-to / recommendation questions ("best way to…", "how do I…", "how to…", "what should I…", "tips for…", "should I…") are no longer classed as external facts. They're answered directly from the model's own knowledge, offline, like any other chat. Genuine facts ("what is the capital of Japan", "how old is …", "who won …") still verify as before. - **v2.1.26** — a *relevance gate* on web results: when a genuine fact is searched and the search returns topically-unrelated pages (the "Best Buy" case), ELI no longer synthesises them into a confident non-answer. It says it searched but couldn't find anything that answers the question, and asks you to rephrase or add a detail — an honest miss instead of a made-up one. Tune with `ELI_WEB_RELEVANCE_FLOOR` (default `0.34`; higher = stricter).

Natural voice / cloning: `pip install "eli-v2.0[natural]"` fails with "No matching distribution"

Symptom: following an in-app message or older doc literally (`pip install "eli-v2.0[natural]"`) fails — pip tries to fetch `eli-v2.0` from PyPI, where this project isn't published. **Cause:** that string only resolves once `eli-v2.0` is already installed under that name from a real index; a source checkout needs the local-path form instead. **Fix:** from the ELI project root: `pip install -e ".[natural]"` (or `.[clone]`/`.[voice]` — all three are aliases for the same Coqui XTTS-v2 extra). Fixed at the source in v2.1.29 — every in-app message now shows the correct form.

Natural voice / cloning: `ImportError: cannot import name 'isin_mps_friendly'`

Symptom: after installing the `[natural]`/`[clone]` extra, the first clone/natural-voice attempt raises `ImportError: cannot import name 'isin_mps_friendly' from 'transformers.pytorch_utils'`. **Cause:** `coqui-tts` (latest release, 0.27.5) still calls a `transformers` helper that existed only as an old Apple-MPS `torch.isin()` fallback and was removed in `transformers>=5` — the exact version ELI's own lock file pins for the rest of the stack, so downgrading transformers isn't an option. **Fix:** fixed in v2.1.29 — `eli/perception/tts_xtts.py` restores the function with a lazy compat shim (plain `torch.isin()`, identical behaviour off MPS) before `TTS` is imported. No action needed on an up-to-date checkout.

Natural voice / cloning: hangs or `EOFError` on first use

Symptom: the very first clone or natural-voice synthesis after installing the extra either hangs indefinitely or raises `EOFError: EOF when reading a line`, and the GUI/voice command silently falls back to a normal voice instead. **Cause:** `coqui-tts` gates the first XTTS-v2 model download behind an interactive y/n terminal prompt (agree to the non-commercial Coqui Public Model License, or confirm a paid licence). ELI's GUI has no terminal for that prompt to read from. **Fix:** fixed in v2.1.29 — `tts_xtts._get_model()` sets `COQUI_TOS_AGREED=1` before loading, which is always the correct branch here (a local, personal voice, never redistributed by ELI). No action needed on an up-to-date checkout; using this feature at all implies agreeing to the non-commercial CPML (a commercial deployment needs a separate paid licence directly from Coqui).

Natural voice / cloning: `ModuleNotFoundError: No module named 'torchaudio'` (or `torchcodec`)

Symptom: `xtts_available()` / a clone attempt fails with a missing `torchaudio` or `torchcodec` import, even though `coqui-tts` installed successfully. **Cause:** `coqui-tts` doesn't declare `torch`/`torchaudio`/`torchcodec` as install dependencies (same reasoning as ELI's own `training` extra — a GPU build must match your CUDA, so it can't be hard-pinned), but genuinely needs all three at import time. **Fix:** fixed in v2.1.29

— the `clone/natural/voice` extras now pull in `torch`, `torchaudio`, and `torchcodec` themselves, so a plain `pip install -e ".[natural]"` on a fresh machine resolves everything together. If you already had a bleeding-edge `torch` installed *ahead* of the latest published `torchaudio`/`torchcodec` (e.g. a nightly build), install the matching version manually from `https://download.pytorch.org/whl/<your-cuda-tag>` — same pattern as the main `torch` line in `requirements.txt`.

---

*House rule: when a new error bites, add it here — symptom, cause, and the **exact** commands that fixed it. Future-you will thank present-you.*

A community plugin won't install

ELI refuses at the first stage that fails, and the dialog names it:

Stage	Meaning	What to do
`listing`	the registry entry is malformed	the publisher must fix their manifest
`payment`	paid plugin, no licence key stored	buy from the publisher, then enter the key
`download`	source URL unreachable	check the URL, and that networking is on
`integrity`	**checksum or signature mismatch**	do not retry — the file does not match what the listing described
`code`	uses a capability it never declared	the publisher must declare it; do not force this
`malware`	a scanner found malicious indicators	do not install
`dependencies`	wants PyPI packages	approve separately with the pip-inclusive button
`consent`	nothing was available to ask you	install from the GUI, not headless

An `integrity` or `malware` refusal is not a bug to work around. Both mean the file on the server is not the file the listing promised, or is actively hostile.

A plugin is installed but does nothing

Installed plugins arrive **switched off with no permissions granted** — that is deliberate. Enable it in Settings ▸ Marketplace ▸ Installed. It will then ask, one capability at a time, the first time it needs each one.

If you chose **Never allow** for a capability, the plugin is never asked again and will keep failing silently at that step. Revoke the decision in Settings ▸ Marketplace ▸ Permissions to be asked afresh.

A plugin asked for nothing and was denied anyway

Permission requests **fail closed** when there is nothing to ask: a headless run, the API server, or a scheduled overnight task. That is intended — a plugin must not gain a capability by running at 3am. Run it once from the desktop GUI and choose **Always allow** if you want it to work unattended.

An MCP server was added but has no tools

Run the doctor (Settings ▸ Marketplace, or `eli.plugins.mcp.doctor()`). It reports each configured server with the exact fault:

- **runtime missing** — `npx`/`uvx` is not installed or not on PATH. The message names the install command for your platform.
- **handshake timeout** — the process started but never answered MCP `initialize`. Usually the wrong command, a missing argument, or a required environment variable.

Note that installing already runs a real handshake, so a server that was accepted did answer at least once. There is exactly one config file — `mcp.config_path()` — and ELI reads no other; a previous config from a different host will not be picked up.

"ELI is offline" but a plugin still reached the internet

Expected, and important to understand. `netguard` patches Python's socket layer **inside ELI's process**. An MCP server, a `pip` install, or any other child process has its own network stack and is not affected. ELI's offline switch does not stop them and ELI cannot see what they send. Only install MCP servers you trust with that. See `security.md` §17.

Training says my GPU can't do it

Read the sentence — it names the shortfall and the fix, e.g. *"6.37 GiB free but this run needs ~9.04 GiB."* Options, in order of effort: close what is using VRAM, lower the sequence length in Advanced, install `bitsandbytes` for 4-bit training, or pick a smaller base model. CPU training works but takes hours to days, which is why it is never selected silently.

"Install" opened a review dialog instead of just installing

That is the design. The one-click path proceeds only when there is nothing to decide. It stops, and names the reason, when: the scan is not clean, the listing has no checksum, the plugin wants permissions, it installs PyPI packages, it costs money, or its download URL is plain `http`. Missing optional scanners and an unsigned plugin do **not** stop it — they are reported afterwards.

A plugin is enabled but raises PermissionError

Runtime enforcement is doing its job. `eli/plugins/sandbox.py` blocks an operation whose capability the plugin's manifest never declared — including one made by importing `socket` or `subprocess` directly rather than through ELI's API. The message names the plugin, the operation and the missing capability.

If the plugin genuinely needs it, the publisher must declare it in the manifest; there is deliberately no way for you to grant a capability the plugin never asked for. `ELI_PLUGIN_SANDBOX=0` disables enforcement entirely and should only be used to confirm that this is what is happening.

A registry or download was "refused for safety"

`netguard.safe_fetch` refuses URLs that are not http(s), and hosts that resolve to loopback, private, link-local or reserved addresses — including after a redirect. This stops a listing you do not control from pointing ELI at your own machine or LAN.

If it is your own registry on your own network, add it as a source: `add_registry` detects a local address, records the exception explicitly, and warns you. A public registry that then redirects somewhere private is still refused.

ELI did something I did not ask for, in the middle of a sentence about something else

Symptom: a long message that mentions an app or a command ("…you did not open spotify or play any song just now, you need to stop lying to me…") stops what is playing, pauses a player or opens a settings

panel. Or a “yes” to ELI’s question runs steps taken from its own explanation.

Cause: the last routing stages match on words wherever they stand, and the reader of ELI’s own offers took a sentence of narrative (“Spotify received that stop signal, paused the track, …”) for a step.

Fix (v2.5.11): a command that changes something is carried out only from a sentence that tells ELI to do it. A message is read sentence by sentence: the commands in it are carried out, and the rest is answered. `eli/runtime/command_splitter.py`, `route_contracts.fits`.

ELI said it had done something and it had not

Symptom: “Spotify is opening now.”, “Done. Playing …”, “Scanning `/home/...` (Waiting for system execution output…)”, “Project `X` initialized.” with nothing happening.

Cause: the message was answered by the model alone, and a model asked to do something writes that it did.

Fix (v2.5.11): on a turn in which nothing was carried out, a sentence that reports an act on the machine is removed and the reply ends “Nothing was carried out in this turn.” If you see that line, the request did not reach a command: say it as an instruction on its own (“open spotify”), or tell ELI which app, file or site you mean.

“open nasa website” answered “not installed, install it?”; “open that website” searched for the words

Cause: the target was taken for an application, and a made-up name was offered for install although no package source had it. “that website” was passed to a search as it stood.

Fix (v2.5.11): a site is opened by its address. An address you give is used as it is; a name (“nasa website”, “the bbc’s home page”) is looked for in the conversation and then on the web; “that website”, “the second link”, “open it in the browser” point at addresses already in the conversation. When no address is found ELI says so and opens nothing. An install is offered only when a package source on your machine has the application.

“maximise browser” changed ELI’s own window, or nothing came back

Cause: there was no rule for maximising a named window, the intent model resolved it without the name, and the command then ran on the window in front, which is ELI’s when you type into it.

Fix (v2.5.11): “maximise firefox”, “fullscreen the browser”, “minimise the browser window”, “switch to firefox” keep the window they name. “switch to the futurama tab” raises the window showing that tab; a tab that is not in front in its window cannot be reached from outside the browser, and the reply says so. On a Wayland session without `kdotool` or `wlrctl`, other programs’ windows cannot be moved at all: the reply says that as well.

I gave ELI a folder and it asked me to run `ls` and paste the output

Fix (v2.5.11): a file or folder written as a path in your message is looked at before the model answers: the reply has in front of it how many files there are, of what kinds, where, and the newest by name, and the documents in the folder (PDF, DOCX, ODT, EPUB, text) are read into the document index, newest first. A large folder is read behind the conversation; ask again after a minute for passages from files that were still queued. Scripts and data files are listed, not read: ask for one by path (“read /path/to/run.py”).

A short remark took half a minute before the reply started

Cause: a message under six words was always sent to the intent model first, which needs its whole action catalogue in the prompt (about 2,400 tokens).

Fix (v2.5.11): a short remark with no question, no command verb, no request and no word in common with anything ELI can do (“doing well, good one”) goes straight to conversation.

Normal mode takes two or three minutes to answer even a short remark

Cause: Normal mode makes two passes, a private scratchpad and then the answer, and each one read the whole prompt (persona, rules, retrieved evidence: about 7.5k tokens). On a hybrid model such as Qwen3.6-35B-A3B llama.cpp cannot reuse a shared start of a prompt, so the second pass paid the full read again (about 65 s on an RTX 2060 SUPER with the experts in RAM). The two passes were also given slightly different system prompts.

Fix (v2.5.12): both passes get one system prompt, and the model’s state at its end is kept between them, so the second pass reads only its own message: 3.2 s instead of 63.8 s, with the same answer. The intent resolver keeps the state of its action list the same way (2.5 s instead of 22-27 s after the first call of a session). The first pass still reads the whole prompt once. `ELI_PROMPT_STATE_REUSE=0` turns this off.

Voice input never works: `transcription failed: TypeError: open() got an unexpected keyword argument 'metadata_errors'`

Cause: the packaged app’s build installed the newest of each package instead of the locked versions, and took PyAV 18. faster-whisper 1.2.1 opens a recording with PyAV and passes an argument PyAV 18 no longer takes, so every transcription failed. Source installs were not affected: the installers keep to `requirements.lock.txt`.

Fix (v2.5.13): ELI reads the recording itself and hands faster-whisper 16 kHz samples, so faster-whisper never opens a file. The release build now installs under the lock as well (`packaging/lock_constraints.py`), and the packaged app’s `--selftest` checks that a recording can be read for speech input.

The first message after start-up takes minutes and the model loads a second time

Cause: the settings named one model (for example a 7B) and the model picked in the start-up dialog was another. Before the first reply ELI compared the two, called the difference a broken “persona lock”, and loaded the model again.

Fix (v2.5.13): the model in memory is the one ELI answers with; a different file in the settings is not a fault, and the model is loaded again only when none is loaded.

ELI’s voice: `PortAudioError: Error opening OutputStream: Invalid sample rate`

Cause: some sound devices (a USB headset’s ALSA device, for one) take only their own sample rate, and the voice is made at 22,050 Hz. Playback fell back to the system player, without the avatar’s lip movement.

Fix (v2.5.13): when the device refuses the voice's rate, the voice is resampled to the device's own rate and played there.

ELI started `ollama serve` at start-up while a GGUF model was in use

Cause: filling in the list of Ollama models (the start-up dialog, the model menu) started the server first; the executor's chat also fell back to any running Ollama when no GGUF model was loaded, and with no GGUF model found the hardware tuner recommended Ollama.

Fix (v2.5.15): Ollama is used, and started, only when it is the backend you chose (in the start-up dialog, Settings or the toolbar's Ollama menu). Listing models never starts it, nothing falls back to it, and ELI never picks it for you.

A library reported usage to its maker

Cause: onnxruntime 1.31, which runs the voice detection of the speech model, sends usage telemetry to Microsoft from its native code, which the network guard cannot see. On first use it writes a device ID and an upload queue to `~/.cache/Microsoft/DeveloperTools`. Hugging Face Hub, Gradio and OpenTelemetry have their own reporting switches.

Fix (v2.5.14): every one of those switches is turned off when ELI starts (`ORT_DISABLE_TELEMETRY`, `HF_HUB_DISABLE_TELEMETRY`, `GRADIO_ANALYTICS_ENABLED`, `DO_NOT_TRACK`, `OTEL_SDK_DISABLED`), and hosts that only collect telemetry are refused by name in `netguard` even with the Net switch on.

Asked why a reply was slow, ELI gave RAM figures

Cause: "that took over a minute??" missed the last-turn rule and reached an older one that answered with the inference-RAM report. Said again as a correction ("I did not ask about RAM"), the same report came back, word for word.

Fix (v2.5.15, v2.5.16): a remark that the last reply was slow goes to the same answer as "why did that take so long": the audit rows of the recent turns, the last one first, with the seconds each took and how much of that was the model, then where the model runs. It is built from the record, not by the model. RAM questions still get the RAM report and "how can we reduce latency" gets the runtime settings. A correction or complaint about the last reply never runs the action that produced it a second time.

"go deeper" on a story from the news briefing got "that story doesn't exist"

Cause: the router's general "go deeper" rule outranked the rule that re-fetches a story after a briefing, so the follow-up went to chat. There, ELI's previous reply had been cut to 280 characters and the story was past the cut; an oversized prompt was also cut from the front. The model had never seen the story, denied it, then called it its own invention.

Fix (v2.5.15): after a briefing, "go deeper on X" fetches X again. The recent exchange is built newest first within its share of the context window: ELI's last reply is whole, a fetched result (a briefing, a search) stays whole for three replies and is labelled as fetched, and when a prompt is too long the oldest lines go, never the message or the newest turns.

“Nothing was carried out in this turn.” under a reply to a greeting

Cause: the claim check adds that line whenever it drops a sentence saying ELI did something; it did so even when nobody had asked ELI to do anything.

Fix (v2.5.15): the line is added only when the message asked ELI to do something (or nothing else of the reply is left). Otherwise the sentence just goes.

A question about ELI got an invented explanation

Cause: “What the fuck is going on!?” matched none of the patterns for a question about ELI, went to the intent model (45 s), and the model invented a “RAG phase” and a “reflection agent”. Where the turn record was given, a long prompt could still cut it out.

Fix (v2.5.15): a short reaction with no subject of its own, an insult aimed at ELI and a correction (“I did not ask about X”, “I asked why …”) are questions about ELI’s last reply: they skip the intent model and get the turn record, which now sits beside the recent exchange where the overflow trim cannot cut it. An insult is no longer recorded as a stance.

The next message waited behind background work

Cause: the proactive daemon judged the model free 45 s after a reply started. A reply streamed for 78 s, so an insight summary and a news synthesis started during it, and the next message waited behind them.

Fix (v2.5.15): a reply in progress counts as busy however long it takes; background work starts only after it ends, and waits about as long as the last reply took.

The runtime report said the context window was “reduced” to a larger number

Cause: llama.cpp rounds the window up to a whole number of its blocks (14,000 asked became 14,080) and the report called any difference “reduced to fit VRAM”.

Fix (v2.5.15): the report says “rounded up” when the window grew and “reduced to fit memory” only when it shrank.

---

v2 used folders and files other ELI versions use too

Cause: an installed v2 kept its data in `~/.local/share/eli`, `~/.config/eli` and `~/.cache/eli`, a few files had generic fallbacks (`~/.eli/`, `~/.eli_user_id`) or relative `artifacts/` paths, and the launcher script and menu entries had generic names (`eli-run`, `eli-setup.desktop`). v2’s setup also removed any ELI menu entry it judged stale, another version’s included.

Fix (v2.5.15): everything of v2’s is under `ELI_v2`, its launchers are `eli-v2-run` and `eli-v2*.desktop`, and setup and uninstall touch only v2’s own entries. Older generic folders are left as they are.