

# ELI — Finetuning Guide

ELI v2.1.68

August 2026

## Contents

|                                                                           |          |
|---------------------------------------------------------------------------|----------|
| <b>ELI Fine-Tuning Guide — Train a Model the Correct Way (End-to-End)</b> | <b>1</b> |
| 0. The five principles (read before touching anything)                    | 1        |
| 1. Concepts you must understand (so this is rigorous, not ritual)         | 2        |
| 2. Environment setup (once)                                               | 2        |
| 3. Step 0 — Choose your model, then get its base HF weights               | 3        |
| 4. The pipeline (the meticulous part)                                     | 3        |
| 5. Verification — is the fine-tune actually good?                         | 5        |
| 6. Troubleshooting (cause → fix)                                          | 5        |
| 7. How this ties into ELI's runtime (so the loop is closed)               | 5        |
| 8. Quick reference (copy-paste — works for ANY chosen base)               | 5        |
| File map                                                                  | 6        |
| Golden rules, one more time                                               | 6        |

## ELI Fine-Tuning Guide — Train a Model the Correct Way (End-to-End)

A meticulous, do-this-exactly walkthrough for turning a **base model you choose** into an **ELI-native GGUF**, then loading it back into ELI. Covers the *why* as well as the *how*, so it's a guide you can reason from, not cargo-cult.

**Method:** QLoRA (4-bit base + LoRA adapters). **Pipeline:** the scripts in training/ (extract → train → merge → convert → install). **Default base:** Qwen/Qwen3-8B — but the whole guide is written around a single BASE=... you set once (Stage 0), so the *same commands* train **any** model you pick into a GGUF.

Scope: this is **fine-tuning** (adapting a pretrained model), **not** pretraining from scratch (that needs thousands of GPU-hours and is out of reach here — and unnecessary).

Hardware: QLoRA keeps this on a single consumer GPU. The defaults target an ~8 GB card; the `--gpu-mem` / base-size knobs in Stage 2 scale it up or down for whatever you have. No specific GPU is assumed anywhere.

---

## 0. The five principles (read before touching anything)

1. **Persona stays dynamic — never bake it into the weights.** ELI's persona is live (persona\_updater + overlay + the per-turn memory brief). You fine-tune the **stable layer only**: ELI's *voice* / *manner* and its *behavioral contracts* (No-Fake-Actions, grounding, dry register). You do **not** train in user facts, memories, dates, or a fixed persona prompt. The old `models/extract_training_data*.py` dumped `memories[-50:]` in as system prompts — that froze state. The new `extract_eli_dataset.py` deliberately does not.

2. **LoRA, not full fine-tune.** A full fine-tune of an 8B model needs ~60–100 GB VRAM. LoRA trains tiny low-rank adapters (<1% of params) on top of a frozen 4-bit base → fits 8 GB. For a *style/voice* shift (what we want), LoRA is not a compromise — it's the right tool.
  3. **Data quality ≠ data quantity.** A few hundred clean, on-voice examples beat thousands of noisy ones. Garbage in → garbage (or overfit) out.
  4. **Match the chat template to the target model.** Qwen3 uses ChatML. The dataset MUST be formatted with the *target model's* tokenizer template, or the model learns the wrong control tokens and produces garbage. (`extract_eli_dataset.py` does this via `tokenizer.apply_chat_template()`.)
  5. **The base model must have enough context for ELI's brief.** ELI's runtime prompt is ~7–9k tokens. A 4k-context base (phi-3-mini) **cannot** run ELI — it produced literal garbage (`_... ,chknce/you, //`). Qwen3-8B's 40,960 context is why it's the pick. (This is enforced at load by the ctx tuner — see §7.)
- 

## 1. Concepts you must understand (so this is rigorous, not ritual)

- **Base vs Instruct weights:** you fine-tune **HF weights** (a directory of `.safetensors` + `tokenizer`), *not* a GGUF. GGUF is the **inference** format ELI loads; HF is the **training** format. They are not interchangeable — you'll download HF for training and produce a GGUF at the end.
  - **LoRA (Low-Rank Adaptation):** freezes the base weights and learns two small matrices per target layer whose product is added to the original weight.  $r$  = the rank (capacity of the change);  $\alpha$  = a scaling factor (effective LR multiplier  $\approx \alpha/r$ ). Small  $r$  (8) = cheap, good for style; large  $r$  = more capacity, more overfit risk.
  - **QLoRA (4-bit):** the frozen base is quantized to 4-bit (nf4) to fit VRAM; the LoRA adapters train in fp16. Quality loss from 4-bit base is negligible for fine-tuning.
  - **target\_modules:** which weight matrices get adapters. We adapt all attention + MLP projections (`q/k/v/o_proj`, `gate/up/down_proj`) — the standard full coverage.
  - **Quantization (inference):** the merged model is converted to GGUF and quantized — `Q4_K_M` for an 8 GB card (best size/quality balance), `Q5_K_M/Q8_0` if you have headroom.
  - **Catastrophic forgetting / overfitting:** too many steps or too high  $r$ /LR makes the model *memorize* your data and *forget* its general ability (and parrot you robotically). Defend with: low  $r$ , modest LR, **few steps**, and watching the loss (see §4).
  - **n\_ctx\_train:** the base model's trained context window. It sets the ceiling for inference ctx; ELI reads it from GGUF metadata at load (the ctx tuner).
- 

## 2. Environment setup (once)

```
# Use the project venv (CUDA already verified there).
```

```
cd ~/Desktop/ELI_MKXI-main_MAY_NEWEST
```

```
.venv/bin/python -m pip install --upgrade unsloth trl transformers datasets accelerate bitsandbytes hugging
```

```
# llama.cpp – needed to convert HF -> GGUF and quantize (if not already present):
```

```
# git clone https://github.com/ggerganov/llama.cpp && cd llama.cpp && make
```

```
# (note its path; you'll point merge/convert at it)
```

Sanity checks:

```
.venv/bin/python -c "import torch; print('CUDA:', torch.cuda.is_available(), torch.cuda.get_device_name(0))"
```

```
.venv/bin/python -c "import unsloth, trl, transformers, datasets; print('train stack OK')"
```

**Hardware honesty:** QLoRA-training an 8B on 8 GB works but is **tight and slow** (heavy CPU offload). If it crawls or OOMs, drop to **Qwen3-4B** as the base (`--base-model Qwen/Qwen3-4B`) — trains far more comfortably, still 32k+ context, still a strong ELI substrate.

---

### 3. Step 0 — Choose your model, then get its base HF weights

#### 0a. Choose the base (the decision that matters most)

You train **HF weights**, not a GGUF (see §1). Pick a base whose context is **≥16k** (ELI's brief is ~7-9k tokens — a 4k base produces literal garbage). Set it **once** as a shell variable and every command below is identical regardless of which model you chose:

```
export BASE="Qwen/Qwen3-8B"           # default – strong ELI substrate, 40k ctx, fits ~8 GB at Q4
# export BASE="Qwen/Qwen3-4B"         # lighter – trains faster/comfier, still 32k+ ctx
# export BASE="microsoft/phi-4"       # 14B dense (MIT), needs more VRAM/offload
# export BASE="tiiuae/Falcon3-10B-Instruct" # or any HF instruct model with ≥16k ctx
export NAME="eli-$(basename "$BASE" | tr '[:upper:]' '[:lower:]')" # output gguf base name
```

| Want                        | Pick                 | Why                                         |
|-----------------------------|----------------------|---------------------------------------------|
| Best all-round default      | Qwen/Qwen3-8B        | 40k ctx, reasoning-capable, fits 8 GB at Q4 |
| Lightest / fastest train    | Qwen/Qwen3-4B        | comfortable on 8 GB, still ≥32k ctx         |
| Strong reasoning, more VRAM | microsoft/phi-4      | 14B dense, MIT-licensed                     |
| Try something else          | any HF instruct repo | <b>must have n_ctx_train ≥ 16k</b>          |

The GGUF keys in ELI's download catalog (qwen3-8b, phi-4, ...) are *inference* builds. For **training** you need the matching **HF** repo (the BASE above), not the GGUF.

#### 0b. Download the base HF weights (~16 GB for an 8B; the one online step)

```
.venv/bin/huggingface-cli download "$BASE" --local-dir "training/base/$(basename "$BASE")"
export BASE_LOCAL="training/base/$(basename "$BASE")" # use this in the commands below
```

You can also pass the HF id directly (`--base-model "$BASE"`) to let HF cache it instead.

### 4. The pipeline (the meticulous part)

#### Step 1 — Extract the dataset (voice/manner only)

```
.venv/bin/python training/extract_eli_dataset.py \
  --base-model Qwen/Qwen3-8B \
  --out training/datasets/eli_voice.jsonl
```

What it does: reads your stored conversations (artifacts/conversations/\*.json), keeps the USER↔ELI turns, formats each with **Qwen3's chat template**, and writes {"text": ...} JSONL. By default `--system-mode none` and **memory-state injection is excluded** (persona stays dynamic).

**Curate before you train (this is where quality is won):** - Open the JSONL. Remove turns where ELI was *wrong*, *off-voice*, or *fabricating* (you don't want to teach the failures we spent this week fixing). - Keep turns that exemplify the voice you want: dry, direct, honest, no fake actions, grounded. - Aim for **a few hundred good examples**, not thousands of mediocre ones. - Optional: hand-write a dozen "gold" exemplars of perfect ELI behavior (refusing to fake an action, admitting uncertainty, the dark-humor register) and append them — high-signal.

#### Step 2 — Train the LoRA adapter

**Always do a smoke run first** (10 steps) to confirm the loop before committing hours:

```
.venv/bin/python training/train_lora.py \
  --base-model Qwen/Qwen3-8B \
  --dataset training/datasets/eli_voice.jsonl \
  --out training/runs/eli-qwen3-8b-lora \
  --max-seq-len 4096 --max-steps 10
```

Then the real run (`--max-steps 200–400`). Hyperparameters, and how to set them:

| Flag                                | Default        | What it controls                         | How to tune                                                                    |
|-------------------------------------|----------------|------------------------------------------|--------------------------------------------------------------------------------|
| <code>--max-steps</code>            | 300            | total optimizer steps                    | start 200–400; <b>watch loss</b> (below). More $\neq$ better.                  |
| <code>--lora-r</code>               | 8              | adapter capacity                         | 8 for voice; 16 only if voice isn't "sticking". Higher $\rightarrow$ overfit.  |
| <code>--lora-alpha</code>           | 16             | adapter scaling ( $\approx 2 \times r$ ) | keep $\approx 2 \times r$ .                                                    |
| <code>--lr</code>                   | 2e-4           | learning rate                            | 1e-4 (gentle) ... 2e-4 (standard). Too high $\rightarrow$ loss spikes/garbage. |
| <code>--max-seq-len</code>          | 4096           | max example length                       | raise to fit your longest multi-turn example; costs VRAM.                      |
| <code>--batch / --grad-accum</code> | 1 / 8          | effective batch = 8                      | leave at $1 \times 8$ on 8 GB; raise grad-accum for a smoother loss.           |
| <code>--gpu-mem / --cpu-mem</code>  | 5.5GiB / 20GiB | VRAM/RAM split                           | lower <code>--gpu-mem</code> if OOM; the rest spills to CPU (slower).          |

**Reading the loss (your only real gauge):** - It should **decrease then flatten**. A gentle decline to a plateau  $\approx$  healthy. - **Loss  $\rightarrow$  near 0** = memorizing your data (overfit). Stop earlier / lower  $r$ /steps. - **Loss flat/noisy from the start** = LR too low, data too small, or template mismatch. - Save adapter checkpoints; if a later checkpoint sounds *more robotic*, use an earlier one.

Output: a LoRA adapter at `training/runs/eli-qwen3-8b-lora/`.

### Step 3 — Merge $\rightarrow$ convert $\rightarrow$ quantize $\rightarrow$ install

```
.venv/bin/python training/merge_and_convert.py \
  --base-model Qwen/Qwen3-8B \
  --adapter training/runs/eli-qwen3-8b-lora \
  --out-name eli-qwen3-8b --quant q4_k_m
```

Produces and installs `models/eli-qwen3-8b-q4_k_m.gguf`. (Manual `llama.cpp` fallback if Unsloth's GGUF export fails — see `training/README.md`.)

### Step 4 — Load it in ELI

Model menu / Auto Detect  $\rightarrow$  pick the new GGUF. The **ctx tuner** reads its real `n_ctx_train` (40960), right-sizes `ctx` to  $\sim 16k$ , and keeps it on the GPU. The **dynamic persona + memory** are still supplied fresh by the runtime — your fine-tune just made the *substrate* speak ELI.

## 5. Verification — is the fine-tune actually good?

Don't trust "the loss went down." Load it and check **four** things: 1. **Voice:** does it sound like ELI (dry, direct, honest) *without* the heavy persona brief leaning on it? Compare a few prompts vs base Qwen3-8B. 2. **Contracts hold:** ask it to "pause spotify" when it can't, or to analyze a file — it must **not fabricate** an action/result (No-Fake-Actions). Ask a grounded question — it must not confabulate. 3. **Persona still dynamic:** confirm the runtime overlay/memory still steer it (e.g., it picks up a new stated preference this session). If it ignores the live brief and only parrots training data → you overfit; retrain gentler. 4. **No capability regression:** a couple of reasoning/coding prompts should be ~as good as base Qwen3-8B. If markedly worse → overfit; fewer steps / lower  $r$ . Keep a tiny held-out set of prompts and run it after every training run — that's your eval.

---

## 6. Troubleshooting (cause → fix)

| Symptom                                           | Cause                                                              | Fix                                                                                                                  |
|---------------------------------------------------|--------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------|
| <b>Garbage output</b> ( <code>_.../you//</code> ) | <code>ctx &gt; model's n_ctx_train</code> , or wrong chat template | ensure base $\geq 16k$ ctx; confirm extract used the <i>target</i> tokenizer template                                |
| <b>CUDA OOM during training</b>                   | base too big for VRAM                                              | lower <code>--gpu-mem</code> ; smaller base (Qwen3-4B); lower <code>--max-seq-len</code> ; <code>--batch 1</code>    |
| <b>Loss won't drop</b>                            | LR too low / data too small / template mismatch                    | raise <code>--lr</code> to $2e-4$ ; add data; re-check template                                                      |
| <b>Model parrots you, ignores live persona</b>    | overfit                                                            | fewer <code>--max-steps</code> , lower <code>--lora-r</code> , earlier checkpoint                                    |
| <b>Worse at reasoning than base</b>               | overfit / too-high rank                                            | lower $r$ , fewer steps                                                                                              |
| <b>GGUF convert fails</b>                         | Unsloth export quirk                                               | use the manual <code>llama.cpp</code> <code>convert_hf_to_gguf.py</code> + <code>llama-quantize</code> path (README) |
| <b>Fine-tune fabricates actions</b>               | trained on ELI's <i>bad</i> turns                                  | scrub the dataset of fabrication/off-voice turns; add gold no-fake exemplars                                         |

---

## 7. How this ties into ELI's runtime (so the loop is closed)

- **Model-agnostic loader:** drop the produced GGUF in `models/`; ELI finds it, no config.
  - **Ctx tuner (shipped):** reads the model's real `n_ctx_train` from metadata, right-sizes ctx, maximises GPU layers, and *warns* if a model's context is too small for ELI's brief — so a bad-fit base can never silently degrade you again.
  - **Contracts stay in code:** No-Fake-Actions, grounding, fenced-JSON tolerance, empty-`<think>` recovery are **runtime guards** (this week's commits). The fine-tune *reinforces* the voice; the guards *enforce* the behavior. Belt and braces.
- 

## 8. Quick reference (copy-paste — works for ANY chosen base)

```
# choose your model ONCE (Stage 0)
export BASE="Qwen/Qwen3-8B"                                     # ← change this to train a different model
export NAME="eli-$(basename "$BASE" | tr '[:upper:]' '[:lower:]')"
export BASE_LOCAL="training/base/$(basename "$BASE")"
```

```
# one-time setup
.venv/bin/python -m pip install --upgrade unsloth trl transformers datasets accelerate bitsandbytes huggingface-cli
.venv/bin/huggingface-cli download "$BASE" --local-dir "$BASE_LOCAL"

# pipeline (extract → train → merge/convert/install)
.venv/bin/python training/extract_eli_dataset.py --base-model "$BASE_LOCAL" --out training/datasets/eli_voic
.venv/bin/python training/train_lora.py --base-model "$BASE_LOCAL" --dataset training/datasets/eli_voic
.venv/bin/python training/merge_and_convert.py --base-model "$BASE_LOCAL" --adapter "training/runs/$NAME-$EPOCHS"
# then: load models/$NAME-q4_k_m.gguf in ELI (Model menu / Auto Detect)
```

## File map

- training/extract\_eli\_dataset.py — dataset (voice only, persona-dynamic-safe)
- training/train\_lora.py — QLoRA trainer (parameterised, Qwen3-8B default)
- training/merge\_and\_convert.py — merge → GGUF → quantize → install
- training/README.md — quick-start; this doc — the full rationale + correctness
- training/datasets/, training/runs/, training/base/ — data, adapters, base weights

## Golden rules, one more time

1. Fine-tune **voice + contracts**, never the dynamic persona/memory.
2. **Clean** the dataset; quality over quantity.
3. **Smoke-run**, then watch the **loss**; stop before it memorizes.
4. Base model must have the **context** to hold ELI's brief ( $\geq 16k$ ).
5. **Verify on real prompts** (voice + contracts + dynamic-persona + no-regression), not on loss.