

ELI — Full Setup Guide

ELI v2.5.31

November 2026

Contents

ELI v2.0 — The Complete Setup Guide (Plain English) **1**
What you're about to install . . . 1
Part 1 — Getting ELI onto your computer . . . 2
Part 2 — Giving ELI a brain (the model) . . . 3
Part 3 — Starting ELI . . . 4
Part 4 — The commands cheat-sheet . . . 5
Part 5 — Understanding the privacy switches . . . 5
Part 6 — When something goes wrong . . . 6
Part 7 — Removing ELI . . . 6

ELI v2.0 — The Complete Setup Guide (Plain English)

> **Updated for v2.4.84 (September 2026).** Release **v2.4.84** on GitHub; cognition pipeline unified (gradient orchestrator, shared retrieval, S01–S12 tracing). Primary install: **prebuilt installers on GitHub Releases** — `ELI-Setup-<v>.exe` (Windows), the `.dmg` (macOS, Apple Silicon), the `.AppImage` (Linux). The Linux AppImage and Windows installer are CI-launch-tested. Data lives in a per-user `ELI_v2` folder and survives upgrades; `--fresh-start` resets it.

Everything you need to install, set up, and run ELI — written for a person, not a programmer. If you can copy and paste, you can do this.

What you're about to install

ELI is an AI assistant that lives **entirely on your own computer**. There's no account to create, no subscription, and nothing you say to it ever leaves your machine unless you deliberately ask for something from the internet (like the news or a model download).

You install it once, download a "brain" (an AI model file) once, and after that it works completely offline.

What you need before starting:

Requirement	Details
A computer	Linux is the best-tested. Windows and macOS installers exist too.
Python 3.10 or newer	Only for the from-source install; the release packages bring their own. Most systems already have it (`python3 --version`). Any version from 3.10 up works, the newest included; with several installed, the installer measures which one the most packages come ready-made for.
Disk space	~2 GB for ELI itself, plus 2–5 GB for a model (more if you pick a big one)
A GPU (graphics card)	**Optional but recommended.** NVIDIA is best supported. Without one, ELI still works — just slower.
Internet	Only for the install itself and the one-time model download.

Part 1 — Getting ELI onto your computer

The easy way (recommended): download a release package

1. Go to the releases page: **https://github.com/ShadowESC95/ELI_v2.0/releases**
2. Download the newest Linux portable file — it looks like: `ELI_v2-<version>-linux-portable.tar.gz`
3. Open a terminal in your Downloads folder and run:

```
tar -xzf ELI_v2-*-linux-portable.tar.gz
cd ELI_v2-*-linux-portable
chmod +x ELI_Setup.sh
./ELI_Setup.sh
```

`ELI_Setup.sh` is the guided, "just press Enter" path — it walks through every step below automatically and opens the setup wizard at the end. **If you used this, you can skip straight to Part 3.**

There's also an AppImage (`ELI_v2-*-x86_64.AppImage`) — make it executable (`chmod +x`), double-click it, done.

The developer way: from the source code

```
git clone https://github.com/ShadowESC95/ELI_v2.0.git
cd ELI_v2.0
bash install.sh
```

That one command does all of this for you:

1. **Looks at your computer** — CPU, RAM, disk space, and what GPU you have (NVIDIA, AMD, Apple, or none).
2. **Shows you a plan** — what it's going to install and why — and asks permission before touching anything.
3. **Creates a private workspace** (a "virtual environment" in a `.venv` folder) so it never interferes with the rest of your system. If you have more than one Python, it picks the one the most of ELI's packages come ready-made for on your machine (it says which and why); `PYTHON=python3.13 bash install.sh` chooses for it.
4. **Installs the AI engine** built for *your* hardware — CUDA build for NVIDIA, ROCm for AMD, Metal for Mac, or a plain CPU build.

5. **Verifies the GPU actually works** — and warns you loudly if you ended up with the slow CPU version by accident.
6. **Installs the helper tools** ELI uses to control your desktop — media playback, screenshots, clipboard, OCR (screen reading), microphone support.
7. **Offers to download a model** sized to your graphics card.
8. **Fetches the small required extras** — the memory "embedder" (~85 MB, needed for ELI's long-term memory) and the voice files (speech-to-text + a speaking voice).
9. **Sets up fresh, empty databases** — ELI starts knowing nothing about you.

Installer options (only if you want them)

You can add these to the end of `bash install.sh`:

Flag	What it means in plain English
`--yes`	"Don't ask me questions, just use the sensible defaults."
`--cpu-only`	"Ignore my GPU, run on the processor." (slower, but always works)
`--install-cuda`	"I have an NVIDIA card but no CUDA toolkit — install that for me too."
`--model=qwen2.5-7b`	"Download this specific model while you're at it."
`--auto-model`	"Pick the best model for my hardware and download it."
`--no-model`	"Don't download anything big — I'll add a model myself later."
`--latest`	"Use the newest versions of everything instead of the tested, frozen set."
`--skip-torch`	"Skip PyTorch." (saves space; disables self-training features)

Example — completely hands-off install with an automatic model:

```
bash install.sh --yes --auto-model
```

Windows

Double-click **`install.bat`** — it launches the full PowerShell installer, which does the same job as the Linux one (CUDA install, dependency lock, GPU verification, database setup). Options:

```
install.bat            (normal — NVIDIA CUDA install)
install.bat /cpu       (no GPU / force CPU)
install.bat /cuda      (also install the CUDA toolkit if missing)
```

Then launch with **`eli.bat`**.

macOS

Same as Linux — `bash install.sh`. It automatically uses Apple's Metal GPU acceleration. The first time ELI takes a screenshot or moves your mouse, macOS will ask you to allow Screen Recording and Accessibility in System Settings → Privacy — grant them and retry.

Part 2 — Giving ELI a brain (the model)

ELI needs a model file (a `.gguf` file) to think with. If you didn't grab one during install, you have three easy options:

Option A — let ELI pick for you (measures your graphics card, picks what fits):

```
.venv/bin/python -m eli.core.model_download --auto
```

Option B — see the menu and choose yourself:

```
.venv/bin/python -m eli.core.model_download --list      # see what's on offer
.venv/bin/python -m eli.core.model_download --choose    # pick any number of them
```

Model	Size	What you need	Who it's for
Qwen2.5-3B	~1.8 GB	4 GB GPU or CPU	Older machines, laptops
Qwen2.5-7B *(default)*	~4.4 GB	8 GB GPU	The sweet spot for most people
Qwen3-8B	~4.7 GB	8 GB GPU	Adds deeper reasoning
Falcon3-10B	~5.9 GB	12 GB GPU	Stronger answers
Phi-4 (14B)	~8.4 GB	12 GB GPU	High quality, MIT licensed
Qwen3.6-35B-A3B	~20.6 GB	24 GB GPU / big CPU	Enthusiast tier
Falcon-H1-34B	~18.9 GB	24 GB GPU / big CPU	Enthusiast tier

Option C — bring your own. Download any chat/instruct `.gguf` file from the internet (Hugging Face is the usual place) and drop it into the `models/` folder. ELI finds it automatically and adapts itself to fit it into whatever memory you have. It is not locked to any brand of model.

The starter pack: the release page also has a tag called `local-assets-v2.1` containing a small starter model, the memory embedder, and voice files. If you have the GitHub `gh` tool set up, `./RUN_ELI.sh --with-github-assets` (portable package) fetches it all in one go.

Part 3 — Starting ELI

The desktop app (the main event)

```
./scripts/eli_launch.sh
```

or simply:

```
./eli.sh
```

The **first launch shows a setup wizard**: it asks what to call you, lets you pick your model and voice, and explains the network toggle. ELI starts as a blank slate — it knows nothing about you until you talk to it.

After install you'll also have **"ELI v2.0" in your app menu** (the installer adds desktop icons on Linux), so you can start it with a click like any other program.

Talking to it

Type in the chat box, or speak — say the wake word (you can set your own; train it in Settings) and just talk. Ask it *"what can you do?"* and it will list everything — all 230 of its capabilities, generated from what's genuinely wired in.

The phone / tablet view (optional)

ELI has a built-in web page you can open from any device **on your own home network**. The AI still runs on your computer — your phone is just a window onto it.

```
./scripts/eli_serve.sh          # just this computer → http://127.0.0.1:8081/
./scripts/eli_serve.sh --lan    # whole home network → prints a protected link + QR code
```

With `--lan` it automatically creates an access token and prints the exact address to type into your phone's browser. Never expose this to the open internet — it's for your home Wi-Fi only.

You can also run both at once (server in the background, desktop app in front):

```
./scripts/eli_launch.sh both --lan
```

Terminal-only mode (no windows)

```
.venv/bin/python -m eli --headless
```

Commands inside it: `/status`, `/mode`, `/reset`, `/help`, `/quit`.

Part 4 — The commands cheat-sheet

Everything in one place. Run these from the ELI folder.

Install & repair

```
bash install.sh                      # full install (asks first)
bash install.sh --yes --auto-model   # fully automatic install + model
./scripts/eli_setup.sh               # guided 8-step first-time setup (the gentle path)
```

Launch

```
./scripts/eli_launch.sh              # desktop app
./scripts/eli_launch.sh serve --lan  # web app for phone/tablet
./scripts/eli_launch.sh both --lan   # both at once
./eli.sh                             # desktop app (shortcut)
.venv/bin/python -m eli --headless   # text-only terminal mode
```

Models

```
.venv/bin/python -m eli.core.model_download --list     # what can I download?
.venv/bin/python -m eli.core.model_download --auto     # pick one for my hardware
.venv/bin/python -m eli.core.model_download --choose   # let me pick from a menu
.venv/bin/python -m eli.core.model_download --aux      # just the memory embedder
```

Voice files (if they didn't download during install)

```
.venv/bin/python -m eli.runtime.voice_assets
```

Health checks

```
bash run_tests.sh                    # run the test suite
bash eli_diagnose.sh                 # system diagnosis report
```

Part 5 — Understanding the privacy switches

These are the promises the software makes, in plain terms:

- **Offline by default.** ELI blocks its own network access at a deep level (the socket layer). When it's offline, its own code *cannot* dial out even if it tries. You flip network access on with one toggle — and ELI announces when it goes online.
- **Your data stays put.** Conversations, your profile, memories — all in local database files inside the ELI folder (`artifacts/`). Delete them any time; ELI simply starts fresh.
- **No account, no telemetry.** Nothing phones home. There's nothing to phone home *to*.

- **One honest caveat:** ELI can run commands and read files on your computer — that's its job as an assistant. Its safety rails are good but not a prison; treat it like a capable tool you're in charge of, not a sandboxed toy. The details live in `SECURITY.md`.

Part 6 — When something goes wrong

Symptom	The fix
"No Python 3 found" or "No Python 3.10 or newer was found"	Install Python from python.org (Windows) or your package manager (`sudo apt install python3`), any version from 3.10 up. Run the installer again; it finds it by itself.
ELI stopped starting after a system or Python upgrade	The workspace was built for the Python you had before. ELI says so when you start it and, if that Python is still installed somewhere, points the workspace back at it with nothing reinstalled. Otherwise run `bash install.sh` (`install.bat` on Windows) to rebuild it: your models, memory and settings are not touched. The release packages are not affected.
ELI is painfully slow	Your AI engine is probably running on CPU. Re-run `bash install.sh --install-cuda` (NVIDIA). The installer tells you at the end whether GPU offload is on.
`.venv not found — run install.sh first`	You skipped the install, or you're in the wrong folder. `cd` into the ELI folder and run `./scripts/eli_setup.sh` (GUI) or `bash install.sh` (terminal).
No sound / voice doesn't work	ELI uses the Piper that was installed with it, not one that happens to be on your system. Run `.venv/bin/python -m eli.runtime.voice_assets`, and make sure `ffmpeg` and `portaudio` installed (the installer prints the exact command if it couldn't do it itself).
First reply after launch is slow	Normal — the model loads into memory on first use. A big model can take a minute.
Out-of-memory / crashes mid-answer	Your model is too big for your GPU/RAM. Download a smaller one (`--auto` picks a safe size).
Phone can't reach the web app	Use `--lan`, make sure both devices are on the same Wi-Fi, and check your firewall allows port 8081.
macOS won't screenshot / move mouse	System Settings → Privacy & Security → grant Screen Recording + Accessibility, then retry.
Wrong model loaded	Settings tab → pick a different model, or drop a new `.gguf` into `models/`.

Still stuck? Open an issue with what you saw: **https://github.com/ShadowESC95/ELI_v2.0/issues** — a plain "I tried it and this happened" is genuinely useful.

Part 7 — Removing ELI

ELI never spreads itself around your system. To remove it completely:

```
rm -rf /path/to/ELI_v2.0        # the folder is the whole install
```

That's it — models, databases, settings, everything lives inside that one folder. (If you added the app-menu icons, remove them with your desktop's app editor or delete the `.desktop` files from `~/.local/share/applications/`.)

ELI v2.0 — © 2026 Jason Fitzgibbon Bridgeman. Source-available under the PolyForm Noncommercial License 1.0.0 (free to use, modify, and share for any noncommercial purpose; not for commercial use). Questions: jaybridgeman0095@gmail.com