



ELI — Full Setup Guide

ELI v2.3.73

September 2026

Contents

ELI v2.0 — The Complete Setup Guide (Plain English)	1
What you're about to install	1
Part 1 — Getting ELI onto your computer	2
Part 2 — Giving ELI a brain (the model)	3
Part 3 — Starting ELI	4
Part 4 — The commands cheat-sheet	5
Part 5 — Understanding the privacy switches	5
Part 6 — When something goes wrong	6
Part 7 — Removing ELI	6

ELI v2.0 — The Complete Setup Guide (Plain English)

Updated for v2.3.73 (September 2026). Release **v2.3.72** on GitHub; cognition pipeline unified (gradient orchestrator, shared retrieval, S01-S12 tracing). Primary install: **prebuilt installers on GitHub Releases** — ELI-Setup-<v>.exe (Windows), the .dmg (macOS, Apple Silicon), the .AppImage (Linux). The Linux AppImage and Windows installer are CI-launch-tested. Data lives in a per-user ELI_v2 folder and survives upgrades; --fresh-start resets it.

Everything you need to install, set up, and run ELI — written for a person, not a programmer. If you can copy and paste, you can do this.

What you're about to install

ELI is an AI assistant that lives **entirely on your own computer**. There's no account to create, no subscription, and nothing you say to it ever leaves your machine unless you deliberately ask for something from the internet (like the news or a model download).

You install it once, download a "brain" (an AI model file) once, and after that it works completely offline.

What you need before starting:

Requirement	Details
A computer	Linux is the best-tested. Windows and macOS installers exist too.
Python 3.10 or newer	Most modern Linux systems already have it. Check with <code>python3 --version</code>
Disk space	~2 GB for ELI itself, plus 2-5 GB for a model (more if you pick a big one)
A GPU (graphics card)	Optional but recommended. NVIDIA is best supported. Without one, ELI still works — just slower.
Internet	Only for the install itself and the one-time model download.

Part 1 — Getting ELI onto your computer

The easy way (recommended): download a release package

1. Go to the releases page: https://github.com/ShadowESC95/ELI_v2.0/releases
2. Download the newest Linux portable file — it looks like: `ELI_v2-<version>-linux-portable.tar.gz`
3. Open a terminal in your Downloads folder and run:

```
tar -xzf ELI_v2-*linux-portable.tar.gz
cd ELI_v2-*linux-portable
chmod +x ELI_Setup.sh
./ELI_Setup.sh
```

`ELI_Setup.sh` is the guided, “just press Enter” path — it walks through every step below automatically and opens the setup wizard at the end. **If you used this, you can skip straight to Part 3.**

There’s also an AppImage (`ELI_v2-*x86_64.AppImage`) — make it executable (`chmod +x`), double-click it, done.

The developer way: from the source code

```
git clone https://github.com/ShadowESC95/ELI_v2.0.git
cd ELI_v2.0
bash install.sh
```

That one command does all of this for you:

1. **Looks at your computer** — CPU, RAM, disk space, and what GPU you have (NVIDIA, AMD, Apple, or none).
2. **Shows you a plan** — what it’s going to install and why — and asks permission before touching anything.
3. **Creates a private workspace** (a “virtual environment” in a `.venv` folder) so it never interferes with the rest of your system.
4. **Installs the AI engine** built for *your* hardware — CUDA build for NVIDIA, ROCm for AMD, Metal for Mac, or a plain CPU build.
5. **Verifies the GPU actually works** — and warns you loudly if you ended up with the slow CPU version by accident.
6. **Installs the helper tools** ELI uses to control your desktop — media playback, screenshots, clipboard, OCR (screen reading), microphone support.
7. **Offers to download a model** sized to your graphics card.
8. **Fetches the small required extras** — the memory “embedder” (~85 MB, needed for ELI’s long-term memory) and the voice files (speech-to-text + a speaking voice).

9. Sets up fresh, empty databases — ELI starts knowing nothing about you.

Installer options (only if you want them)

You can add these to the end of `bash install.sh`:

Flag	What it means in plain English
<code>--yes</code>	"Don't ask me questions, just use the sensible defaults."
<code>--cpu-only</code>	"Ignore my GPU, run on the processor." (slower, but always works)
<code>--install-cuda</code>	"I have an NVIDIA card but no CUDA toolkit — install that for me too."
<code>--model=qwen2.5-7b</code>	"Download this specific model while you're at it."
<code>--auto-model</code>	"Pick the best model for my hardware and download it."
<code>--no-model</code>	"Don't download anything big — I'll add a model myself later."
<code>--latest</code>	"Use the newest versions of everything instead of the tested, frozen set."
<code>--skip-torch</code>	"Skip PyTorch." (saves space; disables self-training features)

Example — completely hands-off install with an automatic model:

```
bash install.sh --yes --auto-model
```

Windows

Double-click **install.bat** — it launches the full PowerShell installer, which does the same job as the Linux one (CUDA install, dependency lock, GPU verification, database setup). Options:

```
install.bat      (normal — NVIDIA CUDA install)
install.bat /cpu  (no GPU / force CPU)
install.bat /cuda (also install the CUDA toolkit if missing)
```

Then launch with **eli.bat**.

macOS

Same as Linux — `bash install.sh`. It automatically uses Apple's Metal GPU acceleration. The first time ELI takes a screenshot or moves your mouse, macOS will ask you to allow Screen Recording and Accessibility in System Settings → Privacy — grant them and retry.

Part 2 — Giving ELI a brain (the model)

ELI needs a model file (a `.gguf` file) to think with. If you didn't grab one during install, you have three easy options:

Option A — let ELI pick for you (measures your graphics card, picks what fits):

```
.venv/bin/python -m eli.core.model_download --auto
```

Option B — see the menu and choose yourself:

```
.venv/bin/python -m eli.core.model_download --list    # see what's on offer
.venv/bin/python -m eli.core.model_download --choose  # pick any number of them
```

Model	Size	What you need	Who it's for
Qwen2.5-3B	~1.8 GB	4 GB GPU or CPU	Older machines, laptops
Qwen2.5-7B (<i>default</i>)	~4.4 GB	8 GB GPU	The sweet spot for most people
Qwen3-8B	~4.7 GB	8 GB GPU	Adds deeper reasoning
Falcon3-10B	~5.9 GB	12 GB GPU	Stronger answers
Phi-4 (14B)	~8.4 GB	12 GB GPU	High quality, MIT licensed
Qwen3.6-35B-A3B	~20.6 GB	24 GB GPU / big CPU	Enthusiast tier
Falcon-H1-34B	~18.9 GB	24 GB GPU / big CPU	Enthusiast tier

Option C — bring your own. Download any chat/instruct .gguf file from the internet (Hugging Face is the usual place) and drop it into the models/ folder. ELI finds it automatically and adapts itself to fit it into whatever memory you have. It is not locked to any brand of model.

The starter pack: the release page also has a tag called local-assets-v2.1 containing a small starter model, the memory embedder, and voice files. If you have the GitHub gh tool set up, ./RUN_ELI.sh --with-github-assets (portable package) fetches it all in one go.

Part 3 — Starting ELI

The desktop app (the main event)

```
./scripts/eli_launch.sh
```

or simply:

```
./eli.sh
```

The **first launch shows a setup wizard**: it asks what to call you, lets you pick your model and voice, and explains the network toggle. ELI starts as a blank slate — it knows nothing about you until you talk to it.

After install you'll also have **"ELI v2.0" in your app menu** (the installer adds desktop icons on Linux), so you can start it with a click like any other program.

Talking to it

Type in the chat box, or speak — say the wake word (you can set your own; train it in Settings) and just talk. Ask it *"what can you do?"* and it will list everything — all 208 of its capabilities, generated from what's genuinely wired in.

The phone / tablet view (optional)

ELI has a built-in web page you can open from any device **on your own home network**. The AI still runs on your computer — your phone is just a window onto it.

```
./scripts/eli_serve.sh          # just this computer → http://127.0.0.1:8081/
./scripts/eli_serve.sh --lan    # whole home network → prints a protected link + QR code
```

With --lan it automatically creates an access token and prints the exact address to type into your phone's browser. Never expose this to the open internet — it's for your home Wi-Fi only.

You can also run both at once (server in the background, desktop app in front):

```
./scripts/eli_launch.sh both --lan
```

Terminal-only mode (no windows)

```
.venv/bin/python -m eli --headless
```

Commands inside it: /status, /mode, /reset, /help, /quit.

Part 4 — The commands cheat-sheet

Everything in one place. Run these from the ELI folder.

Install & repair

```
bash install.sh                # full install (asks first)
bash install.sh --yes --auto-model # fully automatic install + model
./scripts/eli_setup.sh         # guided 8-step first-time setup (the gentle path)
```

Launch

```
./scripts/eli_launch.sh        # desktop app
./scripts/eli_launch.sh serve --lan # web app for phone/tablet
./scripts/eli_launch.sh both --lan # both at once
./eli.sh                       # desktop app (shortcut)
.venv/bin/python -m eli --headless # text-only terminal mode
```

Models

```
.venv/bin/python -m eli.core.model_download --list # what can I download?
.venv/bin/python -m eli.core.model_download --auto # pick one for my hardware
.venv/bin/python -m eli.core.model_download --choose # let me pick from a menu
.venv/bin/python -m eli.core.model_download --aux # just the memory embedder
```

Voice files (if they didn't download during install)

```
.venv/bin/python -m eli.runtime.voice_assets
```

Health checks

```
bash run_tests.sh                # run the test suite
bash eli_diagnose.sh             # system diagnosis report
```

Part 5 — Understanding the privacy switches

These are the promises the software makes, in plain terms:

- **Offline by default.** ELI blocks its own network access at a deep level (the socket layer). When it's offline, its own code *cannot* dial out even if it tries. You flip network access on with one toggle — and ELI announces when it goes online.
 - **Your data stays put.** Conversations, your profile, memories — all in local database files inside the ELI folder (artifacts/). Delete them any time; ELI simply starts fresh.
 - **No account, no telemetry.** Nothing phones home. There's nothing to phone home *to*.
 - **One honest caveat:** ELI can run commands and read files on your computer — that's its job as an assistant. Its safety rails are good but not a prison; treat it like a capable tool you're in charge of, not a sandboxed toy. The details live in SECURITY.md.
-

Part 6 — When something goes wrong

Symptom	The fix
"Python 3.10+ required"	Install Python from python.org (Windows) or your package manager (<code>sudo apt install python3</code>).
ELI is painfully slow	Your AI engine is probably running on CPU. Re-run <code>bash install.sh --install-cuda</code> (NVIDIA). The installer tells you at the end whether GPU offload is on.
<code>.venv</code> not found – run <code>install.sh</code> first	You skipped the install, or you're in the wrong folder. <code>cd</code> into the ELI folder and run <code>bash install.sh</code> .
No sound / voice doesn't work	Run <code>.venv/bin/python -m eli.runtime.voice_assets</code> , and make sure <code>ffmpeg</code> and <code>portaudio</code> installed (the installer prints the exact command if it couldn't do it itself).
First reply after launch is slow	Normal — the model loads into memory on first use. A big model can take a minute.
Out-of-memory / crashes mid-answer	Your model is too big for your GPU/RAM. Download a smaller one (<code>--auto</code> picks a safe size).
Phone can't reach the web app	Use <code>--lan</code> , make sure both devices are on the same Wi-Fi, and check your firewall allows port 8081.
macOS won't screenshot / move mouse	System Settings → Privacy & Security → grant Screen Recording + Accessibility, then retry.
Wrong model loaded	Settings tab → pick a different model, or drop a new <code>.gguf</code> into <code>models/</code> .

Still stuck? Open an issue with what you saw: https://github.com/ShadowESC95/ELI_v2.0/issues — a plain "I tried it and this happened" is genuinely useful.

Part 7 — Removing ELI

ELI never spreads itself around your system. To remove it completely:

```
rm -rf /path/to/ELI_v2.0      # the folder is the whole install
```

That's it — models, databases, settings, everything lives inside that one folder. (If you added the app-menu icons, remove them with your desktop's app editor or delete the `.desktop` files from `~/.local/share/applications/`.)

ELI v2.0 — © 2026 Jason Fitzgibbon Bridgeman. Source-available under the PolyForm Internal Use License 1.0.0 (free to use and modify on your own machine; not for redistribution). Questions: jaybridgeman0095@gmail.com