

ELI — Gui

ELI v2.3.26

August 2026

Contents

ELI GUI	1
Tab map (current — 13 main tabs)	1
Files	2
Launcher (app.py)	2
Main window (eli_pro_audio_gui_v2_0.py)	2
Labs workspace (labs_tab.py)	3
Other surfaces	4
Honest assessment	4

ELI GUI

eli/gui/ — 25.4k LOC, 22 modules (8 top-level + panels/tabs/docks/widgets). A full native PySide6/PyQt desktop app (with a Qt-binding compat shim), plus a first-boot launcher and a large scientific “Labs” workspace.

Tab map (current — 13 main tabs)

flowchart TD

```
W([ELI Pro window]) --> T{13 main tabs}
T --> Chat[Chat]
T --> Pro[Proactive · 6 sub-tabs]
T --> Img[Images]
T --> QA[Quick Actions]
T --> Scr[Screen]
T --> Fil[Files]
T --> Cod[Coding]
T --> Tsk[Tasks · all background jobs]
T --> Rep[Report Builder]
T --> Exp[Experimental · read-only workbench]
T --> Wld[Eli's World]
T --> Set[Settings · 10 pages: Model · Runtime · Generation · Identity · Audio · Application · ]
T --> Labs[Labs · 11 sub-tabs]
Labs --> L1[Notebook · Memory · Jupyter · Calculator · Physics]
Labs --> L2[File Chat · Workspaces · Sim/IDE]
Labs --> L3[Orchestration · Test & Review · dev/diagnostic tools]
Labs --> L4[Training · LoRA wizard]
Set --> S1[Plugins · installed plugin management]
Set --> S2[Marketplace · browse · scan · permissions · sources]
```

New in 2.3.7

Surface	Where	What it is
Training	Labs sub-tab	four-step LoRA wizard: Hardware → Target → Data → Train
Marketplace	Settings sub-tab	community plugins: browse, verify, scan, install, permissions, sources
Permission dialog	anywhere	Android-style consent: Allow once · Always allow · Not now · Never allow
Eval board / Eval + judge	Labs ► Test & Review	behaviour regression board, run in-process

Files

File	LOC	Role
eli_pro_audio_gui_v2_0.py	12.1k	the main window + most app logic (god-file)
labs_tab.py	5.7k	scientific workspace tab
app.py	817	launcher / first-boot auto-tune / entry main()
panels/startup.py	1305	
panels/settings.py	733	settings dialog incl. Plugins + Marketplace sub-tabs
tabs/training_tab.py	841	Labs ► Training — the LoRA wizard
tabs/marketplace_tab.py	707	Settings ► Marketplace — browse/installed/permissions/sources
panels/permission_dialog.py	223	plugin consent dialog + thread-marshalling bridge
docks/operator_console_dock.py	303	operator console dock
widgets/ollama_model_selector.py	294	optional Ollama model picker
tabs/experimental_tab.py,	small	tabs/docks/widgets + Qt compat
panels/agent_wizard.py,		
docks/proactive_dock.py,		
tabs/eli_world_tab.py,		
qt_compat.py, panels/_qt.py		

Launcher (app.py)

The entry path: `_detect_hardware` (queries **free** VRAM — the display server consumes VRAM before ELI launches, so `free` ≠ `total`), `_auto_tune(model_path, hw)` (picks `n_ctx`/layers/batch), `_pick_model`, `_confirm_params`, config load/save. `main()` either shows the startup model picker (first boot) or delegates to `eli_pro_audio_gui_v2_0.main()`. `--setup` forces the wizard.

Main window (eli_pro_audio_gui_v2_0.py)

A 11.0k-line module holding the window **and** a stack of embedded classes that are really application logic, not just UI: - `CentralMemoryAdapter` — bridges the GUI to the memory subsystem. - `LocalModelManager` (708) — discover/load/swap local GGUF models. - `OllamaModelManager` (1142) — optional Ollama integration (legacy/optional; ELI's stance is 100% local GGUF, so this is a secondary path). - `ExecutorBridge`

(1246) — routes GUI actions into the executor/engine. - `_GUIEngineAdapter` — engine façade for the UI. - UI widgets: `_QABoard` (quick-action card board), `_MiniTelemetryGraph` (live telemetry), `_ZoomableSettingsView`, `_ZoomableImagePreview`, `_FlowLayout`, `_CapabilityList`. - `pyqtSignal/Slot` aliased through `qt_compat.py` so it runs on PyQt **or** PySide.

Labs workspace (`labs_tab.py`)

A 5.7k-line “scientific workspace” tab with **11 sub-tabs**: Notebook, Memory & Conversations, Jupyter, Calculator, Physics constants, File Chat, Workspaces, Sim/IDE, **Orchestration, Test & Review, Training**. (Report Builder was promoted out of Labs to its own main tab; Orchestration + Test & Review were demoted back INTO Labs in the 2026-06-18 advisory as developer/diagnostic tools.) This is the research-bench surface (reflects whatever technical/research work the active user does, surfaced dynamically from their own data).

Labs ▶ Training (`tabs/training_tab.py`)

The GUI half of the LoRA pipeline, which `lora_pipeline.py` had always claimed was driven by “the GUI / scheduled task” — the GUI half was never built, so the human review gate the trainer requires had no interface at all.

Four steps, each reporting what it *found* rather than what it assumes, because the operator is not necessarily on the machine this was written on:

1. **Hardware** — accelerator and vendor, free VRAM, whether 4-bit is available, which Python packages are missing, and which trainable base models exist. States plainly when the machine cannot train rather than starting a job that never ends.
2. **Target** — declare a target against any local HF base (family auto-detected), or delete one. Built-ins are listed and cannot be removed.
3. **Data** — the review queue. Filter by needs-review / flagged / clean, read each exchange, approve/reject/edit, bulk-approve the clean set, then save. Offers to mine candidates from stored conversations when the pool is empty, instead of showing an empty table that reads as “you have no conversations”.
4. **Train** — three recipes (Light / Standard / Deep) with the raw parameters behind an *Advanced* disclosure, a readiness check that runs every gate as a dry-run, then the real run on a **QThread** with live step/loss and a cancel.

Threading: training reaches the GUI through **signals only** — a queued signal is the one marshalling primitive that works from a worker, since `QTimer.singleShot` called off the GUI thread never fires.

Settings ▶ Marketplace (`tabs/marketplace_tab.py`)

Four panes matching the four decisions an operator makes: **Browse, Installed, Permissions, Sources**. See `security.md` for the verification and scanning model — the short version is that the install path is deliberately slow, nothing is written to disk before the operator has seen a scan result, and what is written arrives switched off with no permissions granted.

The consent dialog (`panels/permission_dialog.py`)

Modelled on the Android runtime-permission prompt: asked at the point of use, phrased as what the plugin can do *to you* rather than which API it calls, and answerable with **Allow once · Always allow · Not now · Never allow**. The two refusals are distinct on purpose — without a permanent “Never”, a plugin can nag until the operator clicks the wrong button out of fatigue.

`ConsentBridge` marshals requests from worker threads onto the GUI thread and waits, with a bounded timeout that **denies** if the GUI never answers, keeping the fail-closed rule true even when the UI is what failed.

Other surfaces

- `panels/startup.py` — guided first-boot: detect hardware → pick/download a GGUF (HuggingFace) → tune params.
- `panels/agent_wizard.py` — edit an existing agent’s metadata and persona. Authoring a *new* agent now goes through the AgentSpec model (objective, system prompt, triggers, success criteria) rather than a free-text persona — see `orchestration_and_agents.md`.
- `docks/operator_console_dock.py`, `docks/proactive_dock.py` — operator console
 - proactive-suggestion dock.
- `widgets/ollama_model_selector.py`, `tabs/experimental_tab.py`, `tabs/eli_world_tab.py` — optional/experimental surfaces.

Honest assessment

- **Strong:** this is a genuine, feature-rich desktop product — dockable panels, quick-action board, live telemetry graph, zoomable settings/image preview, local model management, a first-boot wizard, an agent-authoring wizard, and a full scientific workspace. Cross-binding (PyQt/PySide) compat is handled. Most local-LLM projects ship a chat box; this is an application.
- **Weak / watch:**
 1. **God-file #3** — `eli_pro_audio_gui_v2_0.py` (11.0k) mixes UI with core logic (`LocalModelManager`, `ExecutorBridge`, `CentralMemoryAdapter`, `_GUIEngineAdapter`). The model/executor/memory bridges should live outside the window module so the UI isn’t coupled to core internals (and so they’re testable headless). `labs_tab.py` (5.7k) is a second large file.
 2. **Ollama manager** (1.1k LOC) sits oddly against the “100% local GGUF, don’t-care-about-Ollama” stance — it’s an optional/legacy path carrying real weight; candidate for removal or clear quarantine.
 3. UI logic instantiating engine/memory directly makes a clean headless mode harder (there is `eli --headless`, but the GUI module re-implements bridges rather than sharing one service layer).

Marketplace panes and the one-click path

Five panes: **Browse**, **Installed**, **Permissions**, **MCP servers**, **Sources**.

Browse carries two buttons on purpose. **↓ Install** is the one-click path — it downloads to memory, checks the checksum and signature, runs every scanner, and then either installs (when nothing needs deciding) or opens the full review dialog with the reasons named. **Check this plugin** is the same work with the dialog always shown, for someone who wants to read the findings first. Neither writes to disk before the operator has an answer.

MCP servers is its own pane because they install differently: there is no source for ELI to scan, only a command. What is checked is that the runtime exists and that the server answers a real `initialize + tools/list` handshake, and the pane states plainly that ELI’s offline switch cannot contain a separate process. Re-check runs `mcp.doctor()` and reports each server’s exact fault.