

ELI — Installation

ELI v2.1.21

July 2026

Contents

Installation — one-click setup, cross-platform (2026-07-05)	1
Which download should I get? (GitHub Releases)	1
Linux — Arch and other lean distros	3
CUDA toolkit option (<code>--install-cuda / /cuda</code>)	3
What <code>bash install.sh</code> does	3
Three local asset layers (not in git)	4
Flags	4
Files	4
Launch	4

Installation — one-click setup, cross-platform (2026-07-05)

Updated for v2.1.21 (July 2026). The primary install path is now the prebuilt installers on GitHub Releases — the Windows Setup.exe and Linux AppImage are CI-launch-tested; the macOS dmg is built on a Mac, best-effort (not CI-verified) with first-boot GPU (CUDA/Vulkan/Metal) and starter-model offers; data lives in a per-user ELI_v2 folder that survives upgrades. Source installs below remain fully supported.

One command per platform sets up Python deps, the **full SQLite architecture** (blank slate — schema only, no personal data), the **nomic embedder**, and the default voice.

Platform	Command	GPU
Linux	<code>bash install.sh</code> (<code>--install-cuda</code> to also fetch the toolkit)	CUDA
macOS	<code>bash install.sh</code>	Metal (no CUDA)
Windows	<code>install.bat / install.bat</code> <code>/cuda (→ install.ps1)</code>	CUDA (winget toolkit)
Android	<code>bash</code> <code>scripts/install_android.sh</code>	CPU only (headless)
Portable tarball	<code>./INSTALL_ELI.sh</code> then <code>./RUN_ELI.sh</code> <code>--with-github-assets</code>	Same as host

Which download should I get? (GitHub Releases)

Every release ships several downloads so you can trade **download size** against **how much works offline out of the box**. Two words explain them all:

- **lean** = source + launchers only (tiny). The Python dependencies **and** a model are fetched on first run, so the first launch needs internet once.
- **full** = the heavy bits are bundled so the first run is faster / more self-contained. On **Windows**, “full” bundles an **offline dependency wheelhouse** (no pip download). On **Linux**, “full” bundles a **starter model** (Qwen2.5-3B, ~1.9 GB) so no model download.

Download	Size	Platform	What’s inside	First run needs internet for	Best for
ELI-Setup-<v>.exe	1 GB	Windows	Guided one-click installer	deps + model	Most Windows users
ELI_v2-<v>-windows-portable.zip	1.2-1.6 GB	Windows	Prebuilt portable app (no install step)	minimal	Windows, unzip-and-run
ELI_v2-<v>-windows-portable-lean.zip (lean)	1.4 MB	Windows	Source + install.bat (online pip)	deps + model	Windows, smallest download
ELI_v2-<v>-windows-portable-full.zip (full)	1.95 GB	Windows	Source + offline dependency wheelhouse	model only	Windows on locked-down/offline networks
ELI_v2-<v>-macos-appimage.dmg	1.2-1.6 GB	macOS (Apple Silicon)	App bundle	model	Mac users
ELI_v2-<v>-linux-portable-lean.tar.gz (lean)	7 MB	Linux	Source + INSTALL_ELI.sh / RUN_ELI.sh	deps + model	Linux, smallest download
ELI_v2-<v>-linux-portable-full.tar.gz (full)	1.9 GB	Linux	Source + bundled Qwen2.5-3B model	deps only	Linux, works with a model immediately
ELI_v2-<v>-x86_64-AppImage (lean)	47 MB	Linux	Double-click app; first launch self-installs	deps + model	Linux, one-click, small download
ELI_v2-<v>-x86_64-full-AppImage (full)	1.9 GB	Linux	Double-click app + bundled model	deps only	Linux, one-click, model included
eli_v2_0-<v>-py3-none-any.whl	2 MB	any	pip wheel (pip install ...whl[full])	deps + model	Developers / CI / existing venvs
SHA256SUMS.txt	tiny	—	SHA-256 for every file above	—	verifying your download

In plain terms: - Want it to just work and don’t mind a one-time first-run download → **Setup.exe** (Windows), **.dmg** (macOS), or the **AppImage (lean)** (Linux). - On a slow or locked-down network → the **-full** variant for your OS (fewer downloads later). - On **Linux** the *model* is the big download, so the **-full** tar / AppImage bundle it (the Python deps still install via pip on first run — internet needed once). On **Windows** the *deps* are the awkward part, so **windows-portable-full** bundles those (the model still downloads). Neither “full” is 100 % offline; each bundles the piece that’s hardest to fetch. - Verify any download against the checksums: `sha256sum -c SHA256SUMS.txt` (Linux/macOS).

Model size vs. your GPU matters more than which download you pick. Choose a model that fits your VRAM — e.g. **Qwen3-8B** / **Qwen2.5-7B** on an 8 GB GPU. Very large models (30B+) run mostly on CPU and will be slow, and the local **web server** may appear to hang on the first message while such a model loads.

Linux — Arch and other lean distros

The **AppImage** is the easiest path on any distro: it bundles its own **Python 3.11** (so Arch’s system Python 3.14, which has no `llama-cpp-python` wheel, is irrelevant) and, since **v2.1.21**, every Qt xcb library it needs — so it launches out of the box with no extra packages. Download and run it **directly**:

```
U=https://github.com/ShadowESC95/ELI_v2.0/releases/download/v2.1.21
wget "$U/ELI_v2-2.1.21-x86_64.AppImage"
chmod +x ELI_v2-2.1.21-x86_64.AppImage
./ELI_v2-2.1.21-x86_64.AppImage
```

Two fixes worth knowing, both resolved in current builds and verified on a clean Arch VM:

- **“attempt to write a readonly database” (portable build).** If ELI is extracted onto an **NTFS / exFAT / FAT or network** filesystem — e.g. a `~/Downloads` on a dual-boot data drive — SQLite’s WAL journal isn’t supported there and database init fails (only the `user.*` stores; the others in the same folder succeed, which is the tell-tale). **v2.1.19+** detects a WAL-hostile filesystem and falls back to a rollback journal automatically. On older builds, extract onto an **ext4/btrfs** path instead. Check a mount with `findmnt -T . -o TARGET,FSTYPE`.
- **GUI won’t open — “libxcb-cursor0 ... xcb platform plugin”.** Qt 6.5+ needs the whole **xcb-util family**, which a minimal desktop doesn’t ship (Qt’s error is misleading — the real missing library is often `libxcb-icccm.so.4`, not the cursor lib). **v2.1.21+** bundles all of them in the AppImage. On an older build or a **source** install, add them:
 - Arch: `sudo pacman -S xcb-util-cursor xcb-util-wm xcb-util-image xcb-util-keysyms xcb-util-renderutil xcb-util`
 - Debian/Ubuntu: `sudo apt install libxcb-cursor0 libxcb-icccm4 libxcb-image0 libxcb-keysyms1 libxcb-render-util0 libxcb-util1`

Run the AppImage **directly** (as above) rather than with `--appimage-extract-and-run`, which unpacks ~4 GB into `/tmp` and can fail with `libz.so.1: file too short` on a small `tmpfs`. Launching over **SSH** into a running desktop session? Prefix with `export DISPLAY=:0`. And paste long download URLs as a **single line** — a wrapped URL 404s and its tail runs as a stray command.

CUDA toolkit option (`--install-cuda / /cuda`)

For non-technical users with an NVIDIA GPU but no toolkit. Best-effort, never fatal: - **Linux:** tries no-`sudo pip nvidia-cuda-nvcc-cu12` (exposes `nvcc` via `CUDACXX`), then the system package manager (`apt/dnf/pacman`, if `sudo` is available), then prints the manual step — then source-rebuilds `llama-cpp` with `-DGGML_CUDA=on`. - **Windows:** `winget install Nvidia.CUDA`, then rebuilds `llama-cpp`. - **macOS/Android:** N/A (Metal / CPU). The default install already uses prebuilt CUDA wheels (no toolkit needed); the option only matters when those don’t match the user’s CUDA or a source build is required.

What `bash install.sh` does

1. Detects Python (3.10+) and OS; creates `.venv`; upgrades `pip/setuptools/wheel`.
2. Installs **PyTorch** (CUDA 12.1 / CPU / macOS-MPS per flags/OS).
3. Installs **llama-cpp-python** with GPU acceleration (CUDA wheel index / Metal / CPU) — then **verifies** `llama_supports_gpu_offload()` and, if it landed CPU-only, prints the exact CUDA-rebuild command (closes the silent-CPU-wheel trap).
4. Installs the ELI package (`[full]`) + all remaining dependencies from the **frozen requirements.lock.txt** (exact known-good versions; reproducible).
5. Seeds `config/settings.json` from the template (offline-by-default, wizard on) — never overwrites an existing config.
6. Runs `python -m eli.core.init_data` — creates **every** SQLite store/table (`user.sqlite3`, `system_index.sqlite3`, `coding_memory.sqlite3`, `agent.sqlite3`) with **zero personal memories/profile/history** (blank slate). System inventory (installed apps, `$PATH` binaries) is scanned once so “open Firefox” works — that is machine environment data, not user memories.
7. Verifies `import eli`, the GUI entry, and the `eli` console script.

8. Fetches **nomic-embed-text-v1.5.Q4_K_M.gguf** → `models/embeddings/` (~80 MiB, required for memory/RAG) via `python -m eli.core.model_download --aux`.
9. Fetches **voice weights**: Piper **en_US-amy-medium** + faster-whisper **small.en** via `python -m eli.runtime.voice_assets` (idempotent; skipped if already present).
10. Optionally downloads a **chat GGUF** (wizard / `--model=` / `--auto`).
11. Prints how to launch + how to add more models.

Three local asset layers (not in git)

Asset	Path	Size (typical)	When fetched
Embedder (required)	<code>models/embeddings/nomic-embed-text-v1.5.Q4_K_M.gguf</code>	~80 MiB	<code>install.sh</code> / wizard / <code>--aux</code>
Chat model (pick one+)	<code>models/*.gguf</code>	0.6–8+ GiB	Wizard / <code>model_download</code> / asset pack
Voice (default amy)	<code>models/tts/piper/</code> + <code>tts_piper/piper/</code>	~60 MiB Piper + ~464 MiB whisper	<code>voice_assets</code> / asset pack

GitHub Release tag **local-assets-v2.1** mirrors embedder + starter chat GGUFs + cleared Piper voices. Restore: `./RUN_ELI.sh --with-github-assets`. Voices excluded from auto-restore: ryan, lessac, cori — see `models/MODEL_LICENSES.md`.

Flags

- `--cpu-only` — no CUDA (CPU torch + CPU llama-cpp).
- `--latest` — use `requirements.txt` version ranges instead of the frozen lock.
- `--skip-torch` — leave an existing torch in place.
- `--no-model` — skip chat-model offer; embedder still fetched unless you disable network.

Files

- `install.sh` (Linux/macOS), `install.bat` / `install.ps1` (Windows).
- `requirements.lock.txt` — frozen exact versions (excludes torch/llama-cpp, which install via their CUDA indices).
- `pyproject.toml` — package metadata + `[project.scripts]` (`eli` → `eli.gui.app:main`).

Launch

`./eli.sh` or `source .venv/bin/activate && eli`. First run shows the **setup wizard** if no chat GGUF is present. Wizard also verifies embedder + voice and can fetch them.

Chat model: `python -m eli.core.model_download --auto` (or `--list`, or a named model). ELI stays offline by default; downloads are deliberate one-time actions.