

ELI V2.0 • NEW USER INSTALL GUIDE

Getting ELI Running

From "I just downloaded this" to "it's talking to me" — on Linux, Windows, or macOS. No jargon you don't need, no steps you can skip. If you can extract a folder and type one line, you can run ELI.

Version: 2.0.15 • **For:** everyone — first-timers and terminal veterans • **Author:** ShadowESC95

The honest bit up front: best-tested on Linux + an NVIDIA card. Windows, macOS and AMD work, but you may hit a rough edge. If you do, tell me — a bug report beats a compliment.

Before you start (30 seconds of honesty)

ELI runs entirely on your own machine. That's the whole point — nothing goes to a cloud. But it also means your computer does the thinking, so a few things are worth knowing.

Thing	What you need
An OK computer	It'll run on modest hardware, but a discrete GPU (NVIDIA is smoothest) makes it fast. No GPU? There's a CPU-only mode — slower, still works.
A little disk space	The app is small. The AI "brain" (a model file) is a separate, larger download you grab once.
Python 3.10+	The installer checks for it and tells you if it's missing. That's the only prerequisite.
Five minutes	Most of it is the installer doing its thing while you get a coffee.

Privacy, in one line

ELI is offline by default and enforces it at the network level. Your conversations, files and memory never leave your computer. You turn the internet on only when *you* want a specific online thing (news, weather, a web search).

Linux — pick one of two easy paths

Path A • The AppImage (the true one-click)

Download `ELI_v2-2.0.15-x86_64.AppImage` from the [Releases page](#), then:

```
chmod +x ELI_v2-2.0.15-x86_64.AppImage # make it runnable (once)
./ELI_v2-2.0.15-x86_64.AppImage
```

The first run installs ELI quietly to a tucked-away folder and sets itself up; every launch after just opens the app. If your file manager lets you double-click AppImages, you can skip the terminal entirely.

Path B • The portable folder (more visible, still easy)

Download `ELI_v2-2.0.15-linux-portable.tar.gz`, extract it, and run the setup:

```
tar -xzf ELI_v2-2.0.15-linux-portable.tar.gz
cd ELI_v2-2.0.15-linux-portable
chmod +x ELI_Setup.sh
./ELI_Setup.sh
```

`ELI_Setup.sh` is the "do everything for me" button — it builds ELI's private workspace, installs what it needs, offers to fetch the AI model, adds a menu icon and an `eli` command, and opens the app. (There's a whole separate guide, *Which script do I run?*, if you want every launcher explained.)

No NVIDIA graphics card?

Install in CPU mode instead — it's the same, just slower to think:

```
./INSTALL_ELI.sh --cpu-only # then:
./RUN_ELI.sh
```

Windows

Two options, both fine:

- **The Setup.exe** (if a release has one) — download and double-click, like any Windows installer. Easiest for non-technical folks.
- **The portable zip** — download `ELI_v2-2.0.15-windows-portable.zip`, extract it, and double-click `install.bat` inside. Then launch with `eli.bat`.

```
install.bat      :: normal install (uses your GPU if present)
install.bat /cpu :: no graphics card
```

PowerShell needs to be allowed to run the installer — Windows may ask; say yes. If Windows SmartScreen frets about an unsigned app, choose *More info* → *Run anyway* (it's unsigned because I'm one person, not a corporation with a code-signing budget).

macOS

Grab the source or portable build and run the same one-click setup as Linux from Terminal:

```
cd ELI_v2-2.0.15-...
./ELI_Setup.sh
```

On Apple Silicon, ELI uses Metal for acceleration (no CUDA). macOS is coded for and works, but it's the least-travelled path — if something's odd, a bug report genuinely helps.

The one extra thing: give ELI a brain

To keep the download small, ELI ships **without** the large AI model. It won't answer until you add one. You have three easy choices:

- **Let the setup do it.** `ELI_Setup.sh` offers to download a model sized to your hardware. Say yes and you're done.
- **Grab the convenience pack** (model + voices) from the project's releases:

```
gh auth login          # log in to GitHub once
./RUN_ELI.sh --with-github-assets
```

- **Bring your own.** Already have a favourite GGUF model? Drop it into the `models/` folder and ELI finds it and fits it to your machine. (Vision models need their matching `mmproj` file too.)

Talk to ELI from your phone

ELI can serve a web page to your phone or tablet on the same Wi-Fi — the same assistant, just on a smaller screen. From the ELI folder:

```
./scripts/eli_serve.sh --lan --https
```

It prints a web address and shows a **QR code** in the app's *Connect a phone* tab. Scan it and you're in. The `--https` part is what lets your phone use its **microphone** — your phone will show a one-time "connection not private" warning (the security certificate is self-signed and local); tap *Advanced* → *Proceed* and the mic unlocks. It all still runs on your computer; nothing touches the cloud.

When it doesn't go perfectly

What you see	What to do
"Virtual environment not found"	You tried to launch before installing. Run <code>./ELI_Setup.sh</code> (or <code>install.bat</code>) first.
Install chokes on GPU / CUDA bits	Reinstall in CPU mode: <code>./INSTALL_ELI.sh --cpu-only</code> (Linux/mac) or <code>install.bat /cpu</code> (Windows).
ELI opens but won't answer	It has no model yet. <code>./RUN_ELI.sh --with-github-assets</code> , or drop a <code>.gguf</code> into <code>models/</code> .
Phone loads but the mic is dead	You need the HTTPS address — start with <code>--https</code> and accept the one-time cert warning on the phone.
Phone can't reach it at all	

Almost always your computer's firewall. ELI prints the exact "allow this port" command when it starts in LAN mode — run it.

The command `eli` runs an old copy Your shell cached it. Type `hash -r` once.

And when in doubt, **ask ELI itself**: "run a full runtime audit", "what are you running on", "why isn't my mic working". It introspects its real state and tells you the truth rather than guessing.

That's it — go talk to it

Open ELI, and on the very first run it'll ask you a couple of friendly questions — your name, what you do, how you like to be worked with. Answer or skip; either way it starts learning you from there. Type or talk, plain English, no command list to memorise.

For the full tour of everything it can do, see *What ELI Is & Can Do* and the *User Manual*. For every command and flag, the *Commands & Installers* reference. Welcome aboard.

ELI v2.0 — a strictly local, private AI assistant that runs on your machine and answers to you. Built by one person; bug reports welcome at github.com/ShadowESC95/ELI_v2.0/issues.