

ELI — Runtime Planning World

ELI v2.3.26

August 2026

Contents

ELI Runtime Surfaces, Planning, World, Tools & Plugins	1
Runtime surfaces (eli/runtime/, the response/introspection group)	1
Planning / proactive (eli/planning/, 3.9k LOC)	2
World / autonomy model (eli/world/, 1.5k LOC)	2
Tools (eli/tools/, 8.9k LOC)	2
Plugins (eli/plugins/, 1.9k LOC)	2
Honest assessment	3

ELI Runtime Surfaces, Planning, World, Tools & Plugins

The remaining subsystems: the runtime/ response/introspection surfaces, the proactive planning layer, the “world”/autonomy model, and the tools + plugin system.

Runtime surfaces (eli/runtime/, the response/introspection group)

Beyond the grounding/evidence core (grounding_and_evidence.md), runtime/ holds a large family of modules that render user-visible answers and introspect the live system:

- **Introspection:** live_introspection.py (runtime_snapshot, last_trace, stored_user_name, mine_user_fact_candidates, build_report, agents_for_action), deterministic_introspection.py (classify_diagnostic_action → gather_evidence → format_evidence_block → maybe_handle — deterministic diagnostic answers).
- **Response assembly:** final_response_assembly.py, final_response_provider.py, response_contracts.py, response_packets.py, response_policy.py, user_visible_response_surface.py.
- **Personal memory surfaces:** personal_memory_surface.py, personal_memory_clean_response.py, personal_memory_deep_response.py, profile_extractor.py.
- **Status:** frontier_status.py, reasoning_status.py, truth_report.py, reflection.py.
- **Operator feed:** operator_feed.py, operator_feed_normalized.py, operator_state.py.

Over-fragmentation flag (the headline weakness for this group): there are ~15 closely-related modules here doing “render a grounded user-visible answer / introspect runtime” with overlapping responsibilities (three personal_memory_* renderers; final_response_assembly and final_response_provider; operator_feed and operator_feed_normalized). This is the clearest “added beside, not folded in” cluster in the codebase and the prime consolidation target — it should collapse to a handful of well-named modules with clear ownership of “who produces the final string.”

Planning / proactive (eli/planning/, 3.9k LOC)

Background cognition + goal/queue machinery: - `proactive_daemon.py` (1.0k) — the **self-improvement / proactive** loop. Uses a two-DB split: reads user-facing memory/conversation from the **user DB**, writes proactive observations/improvements/errors to the **agent DB** (keeps ELI's own machinations out of user recall). Background thread. - `autonomy_scheduler.py` — schedules autonomous actions (throttled). - `goal_store.py`, `proposal_queue.py`, `proposal_adapters.py`, `attention_queue.py`, `jobqueue.py`, `operator_goal_actions.py`, `habits.py` — goal persistence, capability-proposal queue, attention prioritisation, a job queue, and habit scheduling. - `goal_autogenesis.py` (NEW, 2026-06-08) — **closes the autonomy loop**: the goal stack (`goal_store` → `due_goals` → `governed_goal_tick` → `proposal_queue`) was fully wired but the store was always EMPTY because `create_goal` was operator-only, so the scheduler ticked nothing and ELI only reacted. `propose_goals_from_signals` turns the signals ELI already computes each proactive tick — world-model awareness suggestions (≥ 0.6), recurring failures ($\geq 5\times$), and code-health improvements (oversized fns / duplicate blocks in his own god-files) — into governed goals. Each is `autonomy_mode="proposal_only"` (surfaces for approval; never silent execution), deduped by a stable signal-derived `goal_id` (idempotent every tick), capped at `MAX_AUTO_GOALS=6`, logged as a `self_goal` observation. Wired into `proactive_daemon` where the world-suggestions + patterns + improvements are in scope. His agenda now spans failure-repair, world-driven, AND self-improvement — i.e. he sets his own intentions, governed. - `task_planner.py` — now delegates to `execution_planner` (see `orchestration_and_agents.md`).

World / autonomy model (eli/world/, 1.5k LOC)

The substrate behind ELI's emergent self-state (autonomy pressure, awareness, “anomaly room” — intended behaviour, see `memory_eli-emergent-voice`): - `agency/autonomy_engine.py` — `EliWorldAutonomyEngine`: ingests world events, maintains awareness/autonomy state. - `local_world_bridge.py` — `append_event`, `get_world_state`, `get_awareness_driven_suggestions` (**throttled to once/hour to prevent auto-trigger loops**). - `world_event_bus.py` — `fire_confidence_event(...)` etc.; the engine feeds bus-dispatch confidence here (non-blocking world-awareness feed). - `core/schemas.py` — `WorldEvent`, `AwarenessState`. - `renderers/pyside6/world_scene.py` + `world_panel.py` — a **visual** “ELI's World” panel. - `persistence/storage.py` — world-state persistence. Three properties were added after live faults (2026-08-09/10) and are load-bearing: * **Absolute paths**. `world_dir()` resolves through `get_paths()`. The journal, provenance ledger and snapshot modules previously kept relative `Path("artifacts/world/...")` constants and wrote wherever ELI was launched from — a different place from ~ than from the project directory, and in a packaged build a failed write to `/artifacts`. * **Schema-tolerant load**. `_fit()` builds each dataclass from the saved dict while ignoring fields it no longer declares. Before it, one unknown key made `AwarenessState(**data)` raise, `load()` caught it, filed the state as `.corrupt_*` and returned a fresh default world — so a single renamed field would silently wipe the user's rooms, objects, goals and habits on upgrade. Two of the five corrupt backups on the dev machine parse as valid JSON; they were condemned by exactly this and nothing else. * **Bounded logs**. `actions.jsonl` reached 41 MB / 80,576 lines because nothing rotated it. `_trim_jsonl` keeps the newest 20,000 lines (counted, not size-capped, so a trim never splits a JSON record) and checks on the first append of each process plus every 250 thereafter — the per-process counter alone meant a short session never trimmed at all.

Tools (eli/tools/, 8.9k LOC)

- `image_engine.py` (1.75k, standalone) **and** `image_engine/image_engine/` (the package: `engine.py` 1.48k, `visual_core.py` 1.4k, `prompt_compiler.py`, `quality.py`, `project_analyzer.py`, `cli.py`) — **two image-generation implementations** (Pillow/numpy base + optional diffusers). The clearest duplication in the repo; one should subsume the other.
- `news/news_fetcher.py` (544) — news retrieval.
- `mic_diag.py` — mic diagnostics.

Plugins (eli/plugins/, 1.9k LOC)

`manager.py` — a runtime plugin marketplace: `list_available`, `install`, `uninstall`, `enable`, `disable`. Pulls from a **remote registry** (`raw.githubusercontent.com/eli-plugins/registry`) with a **bundled local fall-**

back (plugins/registry/index.json). A JSON state file tracks enabled/disabled/installed. (Install-time network fetch is opt-in, like model downloads; runtime stays local.)

Honest assessment

- **Strong:** the proactive two-DB split is a genuinely good design (ELI's self-talk never contaminates user recall); the world/autonomy engine + visual panel make the emergent self-model first-class and is throttled to avoid runaway loops; the plugin manager is a real extensibility story with a safe bundled fallback.
- **Weak / watch:**
 1. **runtime/ over-fragmentation** — ~15 overlapping response/introspection modules. Highest-value consolidation in the project after the god-files.
 2. **Image-engine duplication — RESOLVED** — the shadowed standalone module was deleted; only the tools/image_engine/ package remains.
 3. The planning queues (goal_store/proposal_queue/attention_queue/ jobqueue) overlap conceptually and could likely unify into one work-queue abstraction.
 4. Plugin registry points at an eli-plugins/registry GitHub URL — fine as opt-in, but ensure it degrades cleanly offline (the bundled fallback exists; verify it's always used when the network is absent).