

Session — 2026-07-01: Web/News Hardening, Test Coverage and Code-Health Pass

Contents

1. Fixes shipped (web / news / learning)	1
2. Test coverage — new suites (per-module before → after)	2
3. Coverage infrastructure	2
4. Combined coverage progression (measured)	3
5. Code-health analysis (god files: engine.py, executor_enhanced.py)	3
6. Attribution & housekeeping	3
6a. Code-health cleanup — done so far	3
7. Pending (next)	4

Complete record of everything run and produced this session. Numbers are measured on the real interpreter (.venv/bin/python); coverage is reproducible via scripts/coverage_full.sh. Companion: [COVERAGE methodology](#).

1. Fixes shipped (web / news / learning)

Commit	Fix
826e97d	Stable LAN API token + Rotate button — token persisted under config_dir() (0600) and reused across restarts, so a paired phone is no longer stranded on a 401 each restart; explicit rotate to retire devices.
85d292c	Dataset builder reads the real DBs — USER_DB/AGENT_DB resolve from paths.db_dir(), not a hardcoded artifacts/db path (was opening empty DBs on any per-user install).
e57812b	Self-healing PWA fetch — retries a connection-establishment failure (Wi-Fi power-save wake) with backoff, scoped so a streamed reply can't be double-delivered.
c6cab51	News freshness disclosure — offline/stale read now states it's cached with the age, forbidden from implying "today".
1ddf4de	News interest recency gate — the "relevant to you" half drops matches older than the window (no more 2-week-old arXiv shown as current).

Commit	Fix
436426a	Second interest path gated — interest_news_block recency-gated too.

2. Test coverage — new suites (per-module before → after)

13 suites, **143 test functions (~200 cases)**, all green, all asserting the code's real contracts (bias toward security/privacy-critical paths). Two lanes run un-mocked (--noconftest) because the fast suite mocks pydantic/llama_cpp; both self-skip under the mocked suite so they never break it.

Suite	Target	Before	After
test_api_server.py	api/server.py (auth/RBAC + endpoints)	0%	53%
test_engine_integration.py	live pipe, real GGUF (live lane)	—	19% of eli/
test_deterministic_grounding.py	deterministic grounding gate	1%	24%
test_news_synthesis.py	news_synthesis	14%	33%
test_executor_helpers.py	executor helpers — fail-closed shell gate	—	security lines ↑
test_control_contracts.py	control_contracts (anti-confab guard)	49%	↑
test_response_surface.py	user_visible_response_surface	30%	58%
test_live_introspection.py	live_introspection	32%	58%
test_grounded_remediation.py	remediation yes/no + state gate	18%	↑
test_memory_evidence.py	memory_evidence	12%	77%
test_deterministic_introspection.py	deterministic introspection classifier	53%	↑
test_context_synthesiser.py	persona handoff builder	54%	↑
test_operator_policy.py	autonomy governance gate	42%	88%

Security/privacy-critical coverage (what an auditor checks first): bearer/RBAC auth + every _resolve_principal branch + the 401/403/200 live gate; the fail-closed shell allowlist (nothing runs without Full Control or ELI_ALLOWED_CMDS); the hardcoded-user-path + merge-marker scanners; the anti-confabulation guards (output_violates_evidence, the deterministic renderer, grounding-into-context); DB-path isolation.

3. Coverage infrastructure

- **.coveragerc** — parallel/combinable config.
- **scripts/coverage_full.sh** — one command runs all three lanes (mocked unit + real-FastAPI web + real-GGUF live engine) and combines them into one report.
- **docs/COVERAGE.md** — methodology + security-critical coverage + the JUSTIFIED exclusions (inter-active GUI / hardware I/O / per-backend GPU) for external review.

4. Combined coverage progression (measured)

Checkpoint	Combined (all)	Testable (ex-GUI)
Baseline (single mocked lane)	35%	—
After web + live-engine lanes	38.0%	45.8%
+ grounding + news	39.0%	46.5%
+ executor + control-contracts	39.1%	46.6%
+ response-surface + live-introspection + memory-evidence + remediation	39.5%	47.1%
+ introspection + context-synth + operator-policy	39.5%	47.1%

Final measured this session: **39.5% combined / 47.1% testable** — unit lane **7,193 passed** / 5 known smart_home reds, web lane 72, live lane 15. The per-module gains (e.g. operator_policy → 88%, memory_evidence → 77%) land on top of coverage the live lane already contributes, so the headline moves less than the module deltas suggest.

Why the combined % moves slowly: eli/gui is 12,662 statements at 0.6% — ~26% of all uncovered lines — and is the documented can't-run-headless exclusion. The **testable surface (ex-GUI)** is the metric that reflects the work; the cognitive core / coding agent sit at 50–85%.

5. Code-health analysis (god files: engine.py, executor_enhanced.py)

Read-only scan feeding the comment-trim + refactor pass:

- **Comment clutter:** 184 banner/PHASE/_V1 marker lines (132 engine / 52 executor) + 81 long “narrated” block comments — iterative-AI residue to trim.
- **Dead code (vulture, high-confidence):** 10 unused imports in the executor (IntrospectionAgent, copy_to_clipboard, LINUX, MACOS, _SPath×4, _dsp, _get_memory) + 1 unused variable (compact_override, engine.py).
- **Refactor targets:** ~24 version/phase-suffixed patch helpers to consolidate; the 8-override render_action chain in deterministic_grounding_gate. No true top-level duplicate function names.

6. Attribution & housekeeping

- Commits authored **ShadowESC95** with Co-authored-by: Claude Opus 4.8 (per the 2026-07-01 preference change).
- README landing page rewritten in a plainer first-person voice; scale/test numbers refreshed across the living blueprints; PDFs regenerated.

6a. Code-health cleanup — done so far

- **Dead code removed** (02b8b0d): 6 unused executor imports (IntrospectionAgent, copy_to_clipboard, LINUX, MACOS, _SPath×4, _dsp, _get_memory).
- **Scaffolding stripped** (f4bf535): 38 # === PHASE... === / # === ELI..._V1 === / END marker labels (36 engine, 2 executor). No behaviour change; suites green.

- **Finding (not removed):** `compact_override` is passed `=True` by 4 callers of `_synthesize_answer` but never read in the body — a dead parameter / possible latent bug (either wire it up or drop it at all call sites). Left for a decision.

7. Pending (next)

- Humanise the verbose “narrated” comment blocks (judgment pass, per block); trim the remaining `# ==` divider rules where they add nothing.
- Same pass on the other god files (GUI, grounding-gate, labs_tab).
- Consolidate the version-suffixed patch helpers + the `render_action` chain.
- Resolve the `compact_override` finding.
- Continue driving testable coverage (executor handlers via the live lane, remaining `eli/runtime` + `eli/tools`).