

ELI MKXI — Test Coverage Suite and Results (Full Report)

Contents

1. Headline	1
2. Methodology — three lanes	1
3. The suite — new test modules (~230 tests)	2
4. Coverage by subsystem	3
5. Why the low subsystems are low (root-cause analysis)	4
6. The 5 pre-existing reds (none from this work)	5
7. On comparisons (a “closest-in-spirit” project claiming ~89%)	5
8. Reproduce	6

A full test-coverage suite was built for ELI and run end-to-end. This document is the complete record: the suite, the three-lane methodology, the measured results (headline + per-subsystem + a visual chart), an honest analysis of the low-scoring areas, the documented exclusions, and how to reproduce every number.

Reproducible via `scripts/coverage_full.sh`; methodology in [docs/COVERAGE.md](#). Measured on the real interpreter (`.venv/bin/python`), 2026-07-02.


1. Headline

Metric	Value
Testable coverage (GUI excluded)	47.6% — 30,764 / 64,587 statements
Raw coverage (whole tree, GUI included)	39.6%
Tests passing	7,218 (unit) + 72 (web) + 15 (live)
New test modules this effort	15 (~230 tests)
Pre-existing reds (unrelated)	5
Skipped / xfailed (documented)	44 / 2

The only exclusion from the “testable” denominator is `eli/gui` (~12,600 statements) — the interactive desktop window can’t be constructed in headless CI (it blocks on `display/device` init). Everything else, including hardware-adjacent and merely-untested code, stays in the count. Excluding untested-but-testable code to pad the figure would be dishonest and is deliberately avoided.

2. Methodology — three lanes

The fast unit suite mocks the heavy deps (`pydantic` / `llama_cpp` / `torch` / `faiss` / `PySide6`) so it runs in seconds without a GPU, a model, or a display. That means the real web server, a real model, and the GUI can’t be exercised in-process — so coverage is measured across three lanes and combined (coverage combine):



Lanes 2 and 3 self-skip under the mocked suite, so they never destabilise the fast run; they're exercised for real on their own lanes and folded in.

3. The suite — new test modules (~230 tests)

Module	Target	Before	After
test_api_server.py	web server: auth/RBAC + every read endpoint	0%	53%
test_engine_integration.py	live pipeline, real GGUF (live lane)	—	19% of eli/
test_operator_policy.py	autonomy governance gate	42%	88%
test_memory_evidence.py	grounded memory bundle	12%	77%
test_response_surface.py	user-visible response coercion	30%	58%
test_live_introspection.py	action→agents map + state readers	32%	58%
test_perception_parsers.py	equation extractor / CSV profiler	39%/12%	100%/78%
test_news_synthesis.py	synthesis helpers + freshness gate	14%	33%
test_news_fetcher_helpers.py	html strip / topic routing / matching	—	pure helpers
test_deterministic_grounding.py	display action contract	11%	24%
test_deterministic_introspection.py	diagnostics action classifier	—	classifier
test_control_contracts.py	anti-confabulation guard	49%	↑
test_context_synthesiser.py	persona handoff builder	54%	↑

Module	Target	Before	After
test_grounded_remediation.py	yes/no intent + repair state	18%	↑
test_executor_helpers.py	fail-closed shell gate + scanners	—	security lines


Bias throughout: the **security/privacy-critical** paths an auditor checks first — bearer/RBAC auth, the fail-closed shell allowlist, the hardcoded-path/PII scanners, the anti-confabulation guards, and DB-path isolation.

4. Coverage by subsystem

Subsystem	Cover	Statements	Note
eli/onboarding	85%	small	interview flow
eli/coding	81%	~1,045	the frontier coding agent reasoning
eli/cognition	65%	~7,021	/ agents / persona
eli/learning	62%	~1,520	LoRA pipeline
eli/core	54%	~3,492	paths / hardware / netguard
eli/memory	54%	~3,338	SQLite + FAISS + KG
eli/world	54%	~971	world/autonomy model
api (web server)	51%	~1,144	FastAPI + dashboard
eli/kernel	51%	~7,059	the cognition engine
eli/runtime	51%	~13,131	grounding / guards / contracts
eli/contracts	44%	~256	typed contracts
eli/planning	41%	~2,132	daemon / scheduler / autonomy

Subsystem	Cover	Statements	Note
eli/execution	41%	~12,315	router (70%) + executor handlers
eli/plugins	38%	~1,050	runtime-loaded tools
eli/system	37%	~261	system index
eli/tools	~30%	~4,510	news / image / weather
eli/perception	~27%	~4,138	vision / audio / OS / parsers
eli/utils	23%	~468	logging + platform_compat
eli/integrations	21%	tiny	mostly empty/stub
eli/gui	excluded	~12,662	interactive, can't run headless

Visual (testable surface)



5. Why the low subsystems are low (root-cause analysis)

The low scores cluster around **one cause: the I/O boundaries where ELI touches the real world.** Pure logic is well covered; the edges aren't.

Subsystem	Why it's low
execution (41%)	The executor is ~174 action handlers, most side-effecting — open apps, run shell, screenshots, volume, file writes, media. You can't unit-test "open Firefox" (the test <i>does it</i>). They cover only via real live-lane turns (a limited safe set). The router beside it is 70% — pure parsing.
perception (27%)	The body: GPU vision, whisper STT (mic), TTS (speakers), gaze (webcam), os_controller (8% — needs a live desktop), screenshots (a display). None runs headless . The covered part is the pure parsers (equations 100%, CSV 78%).
utils (23%)	Almost entirely platform_compat.py — if WINDOWS ... elif MACOS ... else ... branches. On the Linux CI box only the Linux branch runs ; the other-OS code is counted uncovered. Inherent to cross-platform code.
tools (30%)	News fetcher (network, gated off in tests), image engine (GPU diffusion), weather (network). The non-network logic is tested; fetch/GPU paths aren't.
planning (41%)	Proactive daemon / autonomy tick / scheduler run on background timer threads tests don't drive. The pure scheduling logic (infer_kind, parse_when) is tested.
plugins (38%) / integrations (21%)	Runtime-loaded / mostly-stub. Not all exercised; integrations is near-empty (noise on a tiny denominator).

Closeable vs not: - *Closed this pass:* the pure parsers and news helpers (equations, CSV, HTML/topic routing) — moved from ~10–39% to 78–100%. - *Still closeable:* more file parsers, tools/planning logic, and **more live-lane turns** for the executor. This is the honest path past 47%. - *Not honestly closeable in CI:* GPU-vision, mic/camera/speaker, live-desktop control, and other-OS branches — need real hardware or another OS. Forcing them "green" with mocks would only test the mocks.

6. The 5 pre-existing reds (none from this work)

1-3. smart_home plugin — the in-progress Home-Assistant removal (voice SMART_HOME now uses ELI's own MQTT server). 4. A blueprint references a since-moved file (eli/execution/handlers/__init__.py). 5. Silent-swallow ratchet — 987 except: pass vs a 950 ceiling (a standing observability debt; the ratchet test correctly forbids raising the ceiling — clear it by making the swallows observable, not by hiding them).

7. On comparisons (a "closest-in-spirit" project claiming ~89%)

The gap is largely *what ELI is*. ELI carries a large **untestable embodiment surface** — desktop GUI, gaze/webcam, mic/voice, local vision, OS control, smart-home — that can't run in headless CI. A leaner pure-software agent simply lacks that surface, so a much higher fraction of its code is unit-testable by construction. ELI's cognitive **core is already in a comparable band** (coding 81%, cognition 65%, memory 54%); the raw number is lower because ELI does more, not because the core is weak. The honest way past 47% is more **live-lane integration tests**, not denominator tricks.

8. Reproduce

```
bash scripts/coverage_full.sh          # runs all 3 lanes, combines, reports  
# outputs: artifacts/coverage_report.txt + artifacts/coverage_html/index.html
```

Config: .coveragerc (GUI-only omit). The exclusion is justified in docs/COVERAGE.md. Numbers in this report are regenerated by that command on any checkout.